Vertex AI AutoML

1. Movie Rating prediction

Trained Model image

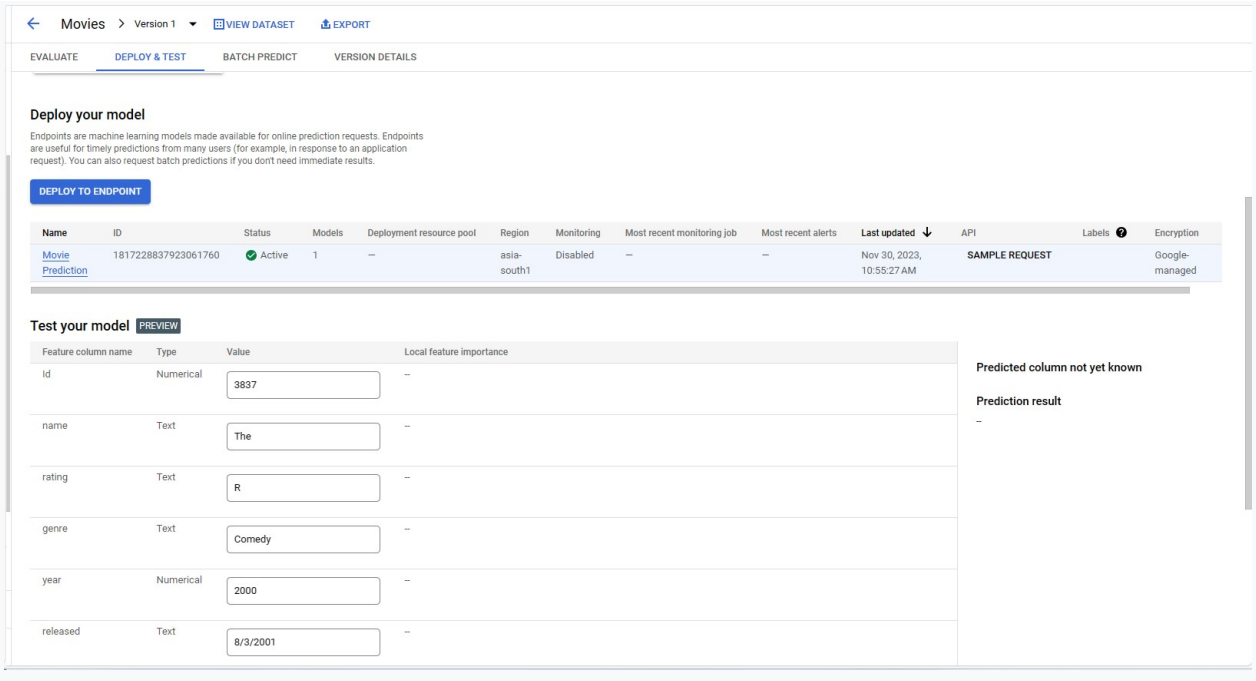

Trained model deployment

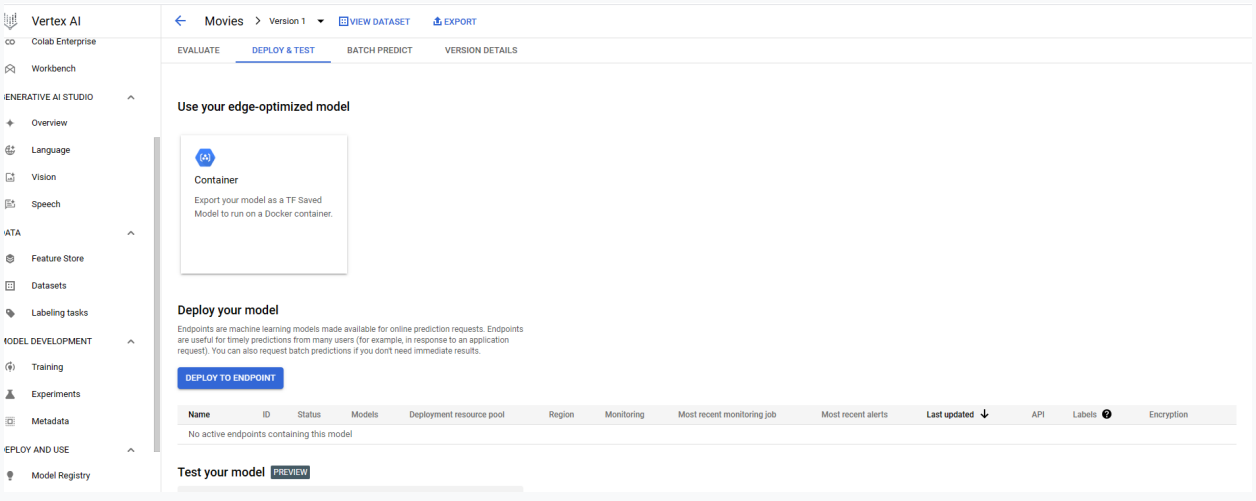

Trained Model Test/Prediction image

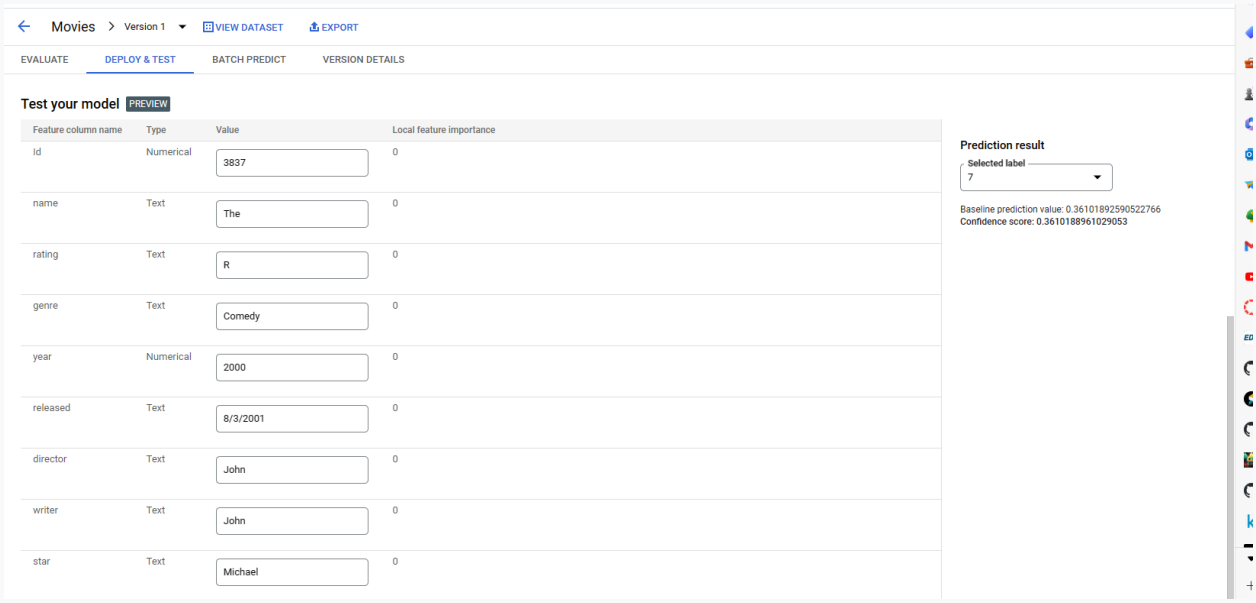

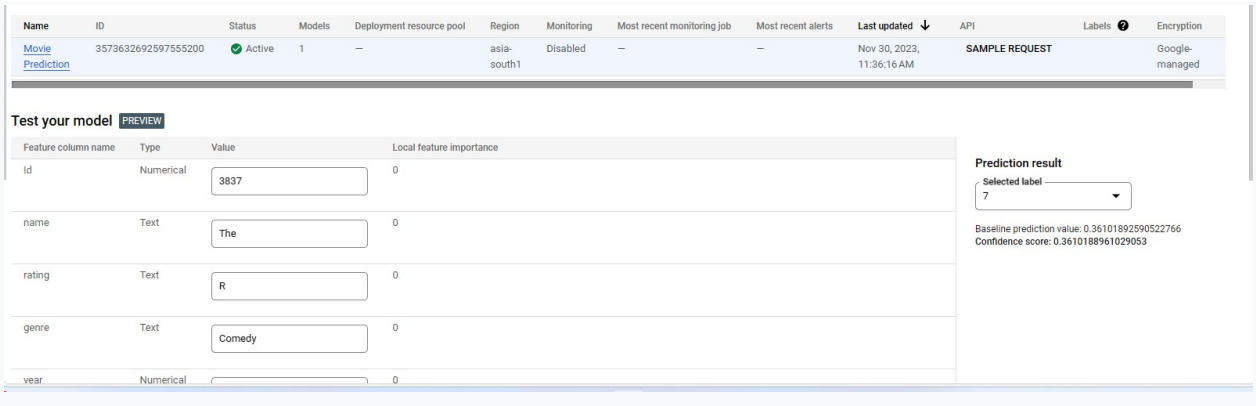

Evaluation Details

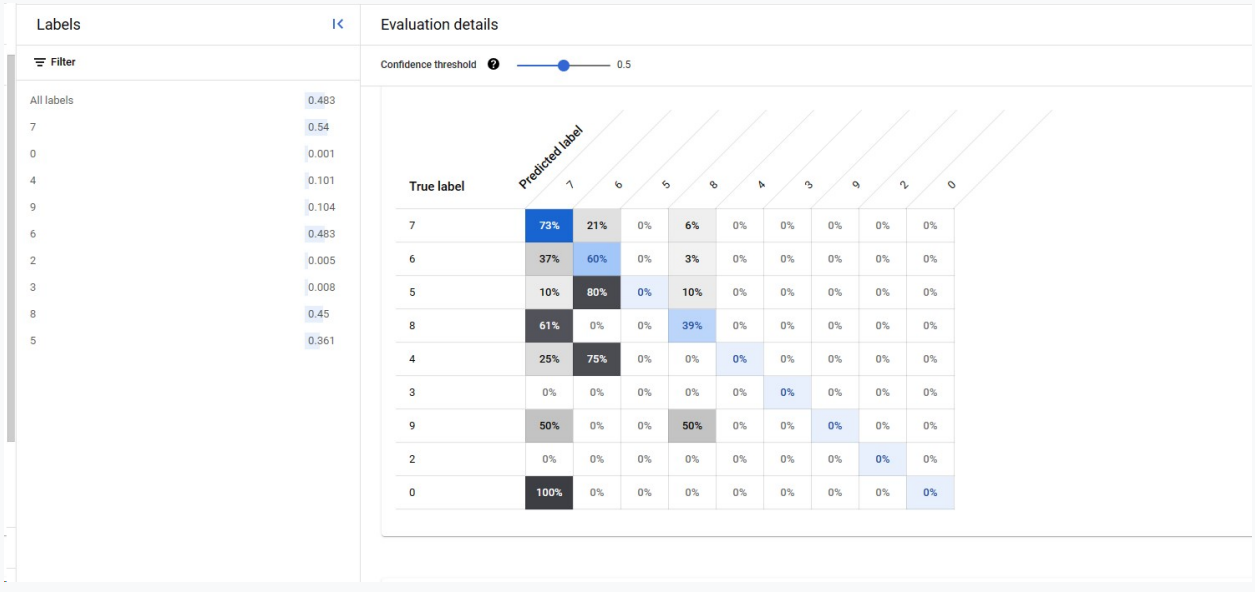

ALL other images:

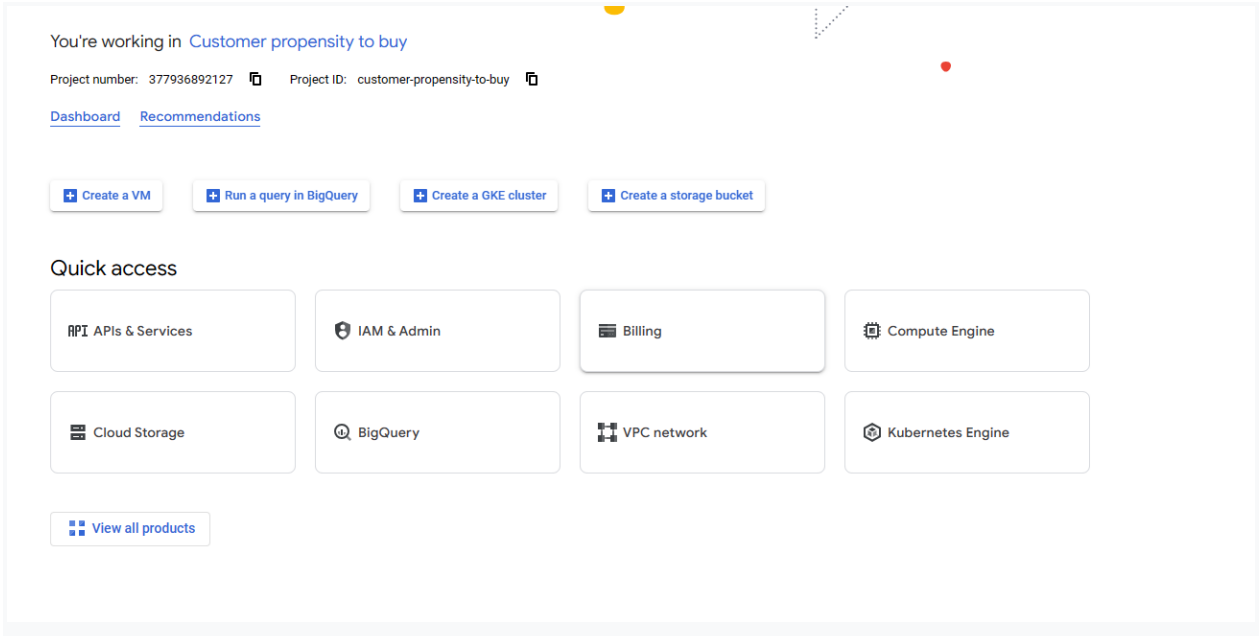
You're working in Customer propensity to buy
Project number: 377936892127
Project ID: customer-propensity-to-buy
Dashboard
Recommendations
Create a VM
Run a query in BigQuery
Create a GKE cluster
Create a storage bucket
Quick access
APIs & Services
IAM & Admin
Billing
Compute Engine
Cloud Storage
BigQuery
VPC network
Kubernetes Engine
View all products

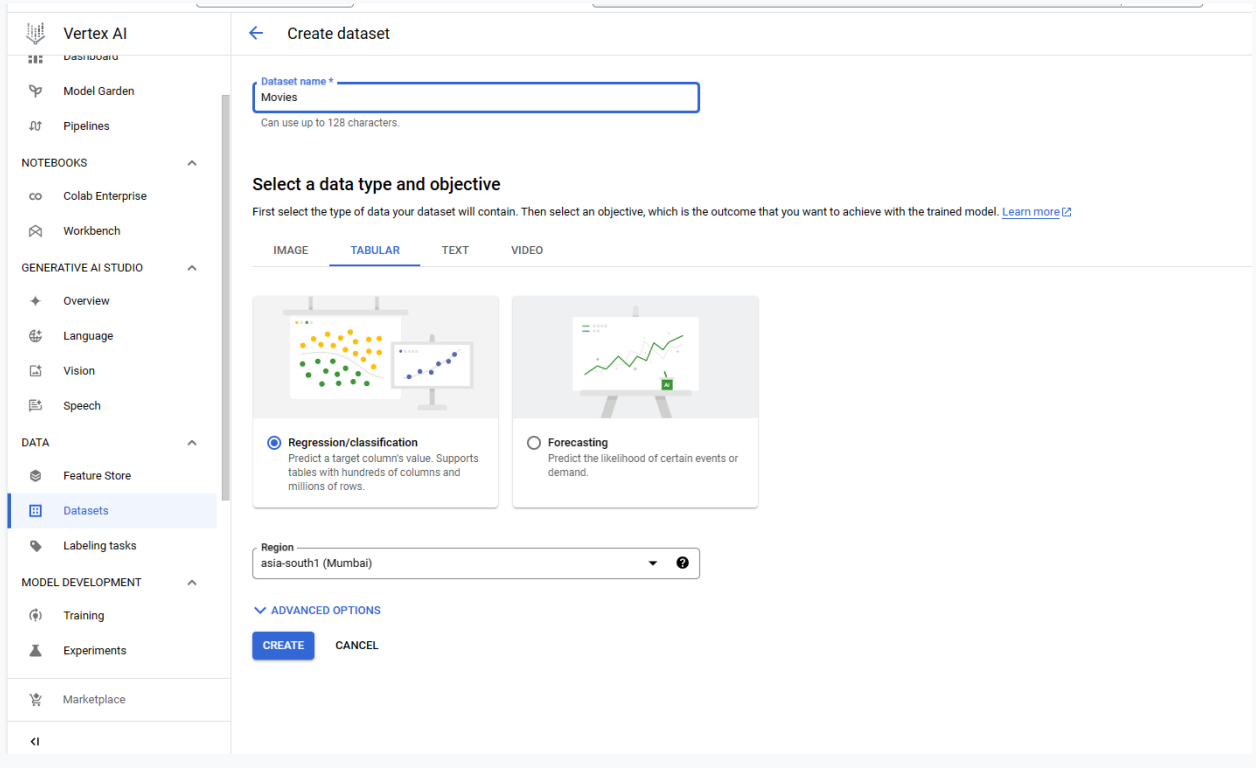
Vertex AI
Model Garden
Pipelines
NOTEBOOKS
Colab Enterprise
Workbench
GENERATIVE AI STUDIO
Overview
Language
Vision
Speech
DATA
Feature Store
Datasets
Labeling tasks
MODEL DEVELOPMENT
Training
Experiments
Marketplace
Create dataset
Dataset name *
Movies
Can use up to 128 characters.
Select a data type and objective
First select the type of data your dataset will contain. Then select an objective, which is the outcome that you want to achieve with the trained model. Learn more
IMAGE
TABULAR
TEXT
VIDEO
Regression/classification
Predict a target column's value. Supports tables with hundreds of columns and millions of rows.
Forecasting
Predict the likelihood of certain events or demand.
Region
asia-south1 (Mumbai)
ADVANCED OPTIONS
CREATE
CANCEL

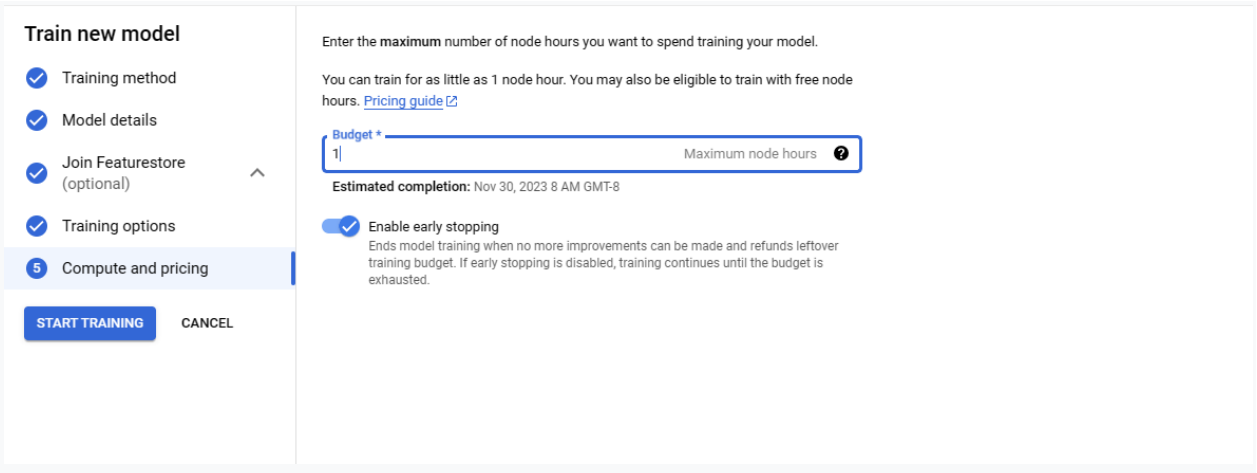
Train new model
Training method
Model details
Join Featurestore (optional)
Training options
5
Compute and pricing
START TRAINING
CANCEL
Enter the maximum number of node hours you want to spend training your model.
You can train for as little as 1 node hour. You may also be eligible to train with free node hours. Pricing guide
Budget *
1
Maximum node hours
Estimated completion: Nov 30, 2023 8 AM GMT-8
Enable early stopping
Ends model training when no more improvements can be made and refunds leftover training budget. If early stopping is disabled, training continues until the budget is exhausted.

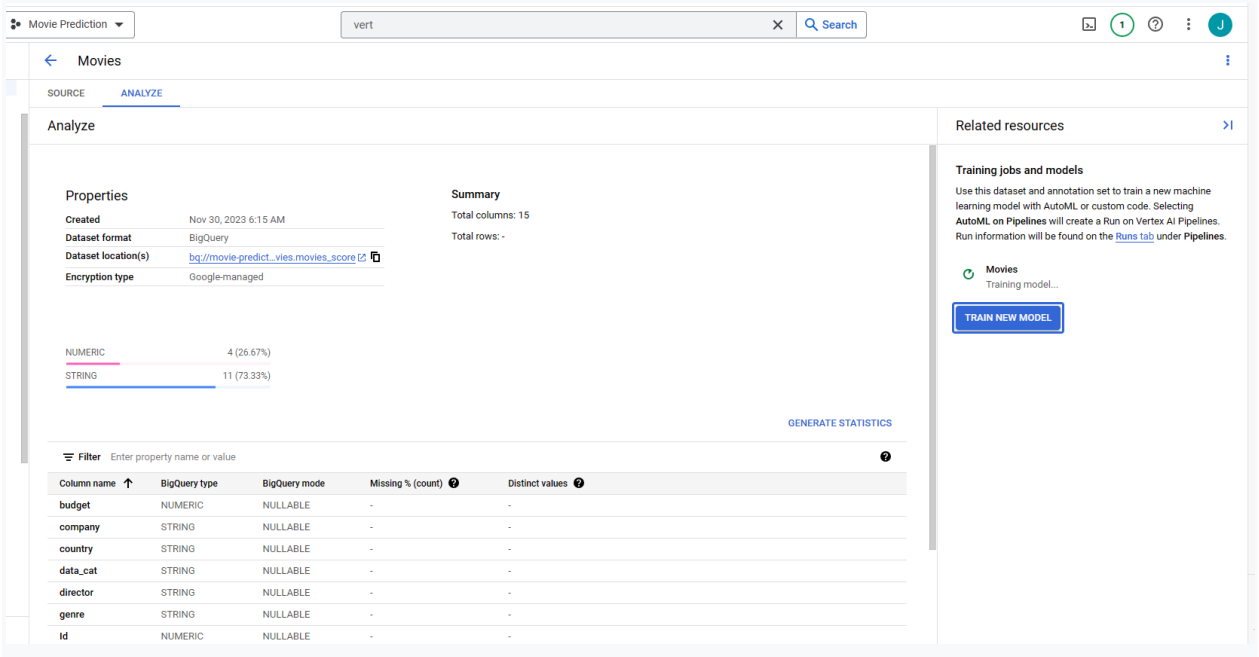

Training began at Nov 30, 2023, 6:43:41 AM and is still in progress.

Status	Training
Training pipeline ID	7909704131290333184
Created	Nov 30, 2023, 6:42:41 AM
Start time	Nov 30, 2023, 6:43:41 AM
Budget (original)	1 node hours
Elapsed time	1 hr 58 min
Region	asia-south1
Encryption type	Google-managed

Dataset	Movies
Target column	score
Data split	Randomly assigned (80/10/10)
Column metadata	VIEW DETAILS

Algorithm	AutoML
Objective	Tabular classification
Optimized for	Log loss
Training stage	Model post processing

Training performance

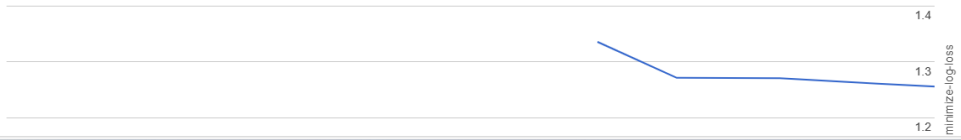

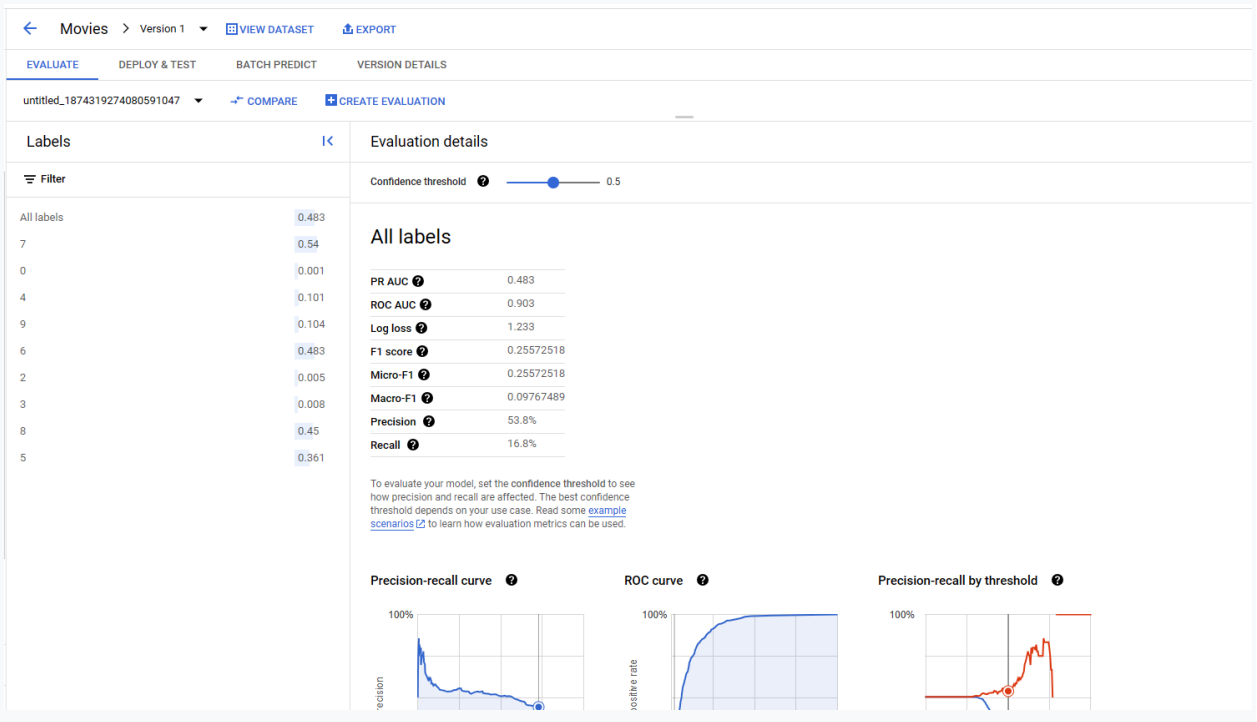
Movies
Version 1
VIEW DATASET
EXPORT
EVALUATE
DEPLOY & TEST
BATCH PREDICT
VERSION DETAILS
untitled_1874319274080591047
COMPARE
CREATE EVALUATION
Labels
Filter
All labels 0.483
7 0.54
0 0.001
4 0.101
9 0.104
6 0.483
2 0.005
3 0.008
8 0.45
5 0.361
Evaluation details
Confidence threshold 0.5
All labels
PR AUC 0.483
ROC AUC 0.903
Log loss 1.233
F1 score 0.25572518
Micro-F1 0.25572518
Macro-F1 0.09767489
Precision 53.8%
Recall 16.8%
To evaluate your model, set the confidence threshold to see how precision and recall are affected. The best confidence threshold depends on your use case. Read some example scenarios to learn how evaluation metrics can be used.
Precision-recall curve
ROC curve
Precision-recall by threshold
100%
100%
100%

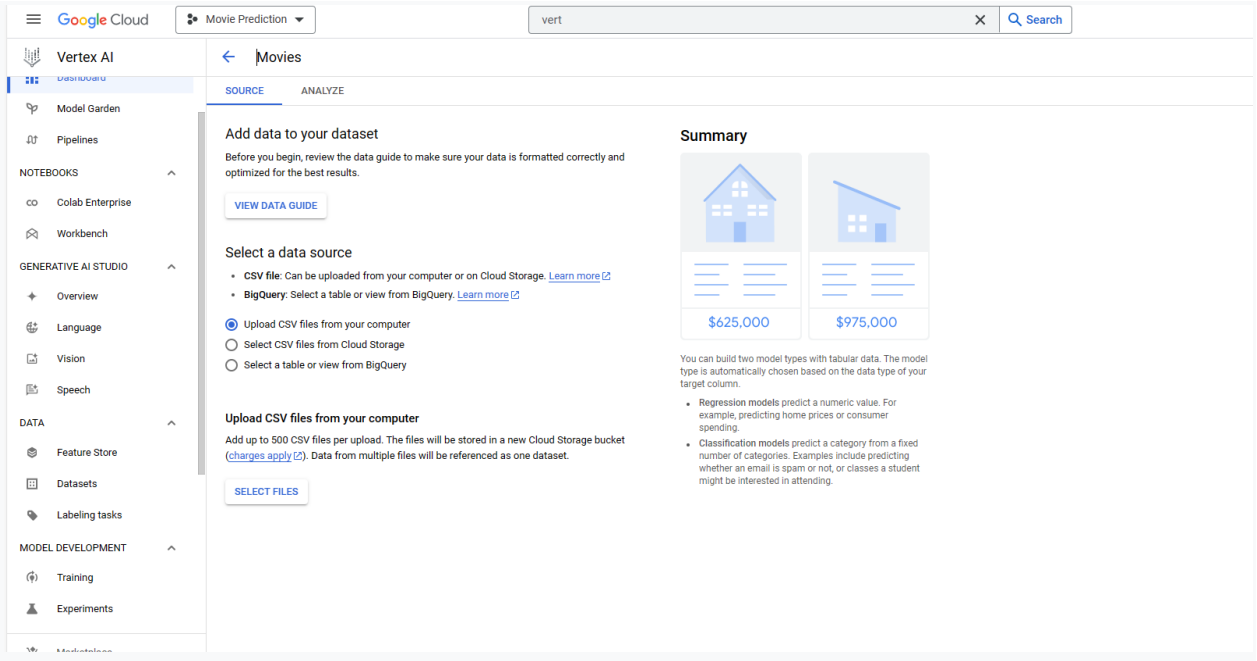
Google Cloud
Movie Prediction
vert
Search
Vertex AI
Model Garden
Pipelines
NOTEBOOKS
Colab Enterprise
Workbench
GENERATIVE AI STUDIO
Overview
Language
Vision
Speech
DATA
Feature Store
Datasets
Labeling tasks
MODEL DEVELOPMENT
Training
Experiments
Movies
SOURCE
ANALYZE
Add data to your dataset
Before you begin, review the data guide to make sure your data is formatted correctly and optimized for the best results.
VIEW DATA GUIDE
Select a data source
CSV file: Can be uploaded from your computer or on Cloud Storage. Learn more
BigQuery: Select a table or view from BigQuery. Learn more
Upload CSV files from your computer
Select CSV files from Cloud Storage
Select a table or view from BigQuery
Upload CSV files from your computer
Add up to 500 CSV files per upload. The files will be stored in a new Cloud Storage bucket (charges apply). Data from multiple files will be referenced as one dataset.
SELECT FILES
Summary
$625,000
$975,000
You can build two model types with tabular data. The model type is automatically chosen based on the data type of your target column.
Regression models predict a numeric value. For example, predicting home prices or consumer spending.
Classification models predict a category from a fixed number of categories. Examples include predicting whether an email is spam or not, or classes a student might be interested in attending.

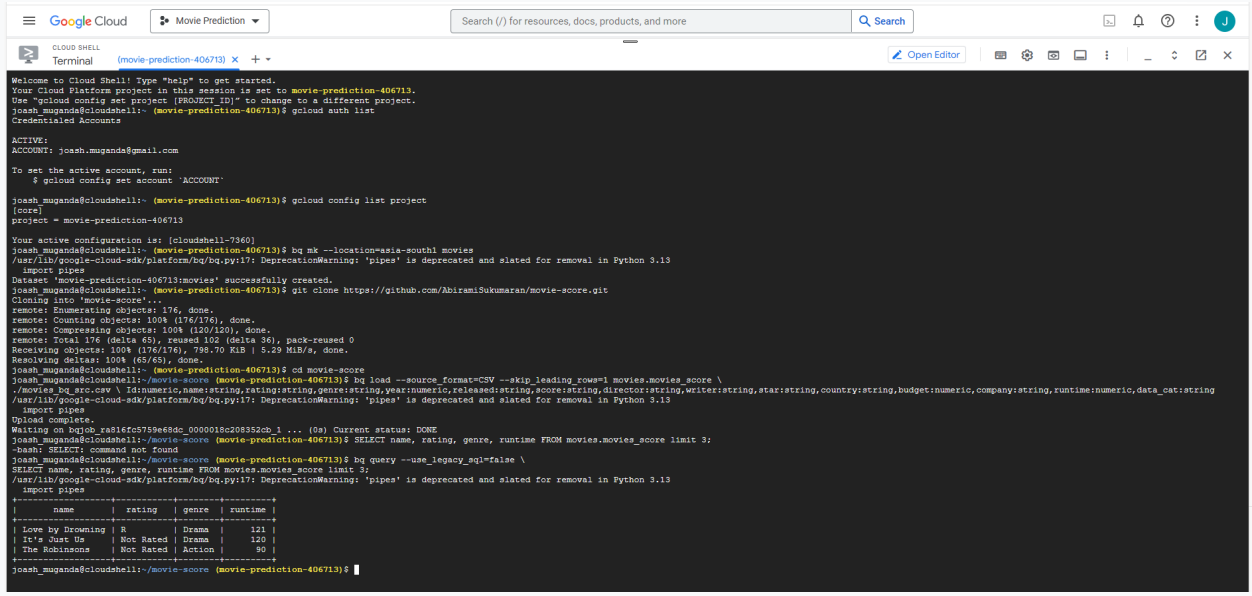
Google Cloud
Movie Prediction
Search (/) for resources, docs, products, and more
Search
CLOUD SHELL
Terminal
(movie-prediction-406713)
Open Editor
Welcome to Cloud Shell! Type "help" to get started.
Your Cloud Platform project in this session is set to movie-prediction-406713.
Use "gcloud config set project [PROJECT_ID]" to change to a different project.
joash_muganda@cloudshell:~ (movie-prediction-406713)$ gcloud auth list
Credentialed Accounts
ACTIVE:
ACCOUNT: joash.muganda@gmail.com
To set the active account, run:
$ gcloud config set account `ACCOUNT`
joash_muganda@cloudshell:~ (movie-prediction-406713)$ gcloud config list project
[core]
project = movie-prediction-406713
Your active configuration is: [cloudshell-7360]
joash_muganda@cloudshell:~ (movie-prediction-406713)$ bq mk --location=asia-south1 movies
/usr/lib/google-cloud-sdk/platform/bq/bq.py:17: DeprecationWarning: 'pipes' is deprecated and slated for removal in Python 3.13
import pipes
Dataset 'movie-prediction-406713:movies' successfully created.
joash_muganda@cloudshell:~ (movie-prediction-406713)$ git clone https://github.com/AbiramiSukumaran/movie-score.git
Cloning into 'movie-score'...
remote: Enumerating objects: 176, done.
remote: Counting objects: 100% (176/176), done.
remote: Compressing objects: 100% (120/120), done.
remote: Total 176 (delta 65), reused 102 (delta 36), pack-reused 0
Receiving objects: 100% (176/176), 798.70 KiB | 5.29 MiB/s, done.
Resolving deltas: 100% (65/65), done.
joash_muganda@cloudshell:~ (movie-prediction-406713)$ cd movie-score
joash_muganda@cloudshell:~/movie-score (movie-prediction-406713)$ bq load --source_format=CSV --skip_leading_rows=1 movies.movies_score \
./movies_bq_src.csv \ Id:numeric,name:string,rating:string,genre:string,year:numeric,released:string,score:string,director:string,writer:string,star:string,country:string,budget:numeric,company:string,runtime:numeric,data_cat:string
/usr/lib/google-cloud-sdk/platform/bq/bq.py:17: DeprecationWarning: 'pipes' is deprecated and slated for removal in Python 3.13
import pipes
Upload complete.
Waiting on bqjob_ra816fc5759e68dc_0000018c208352cb_1 ... (0s) Current status: DONE
joash_muganda@cloudshell:~/movie-score (movie-prediction-406713)$ SELECT name, rating, genre, runtime FROM movies.movies_score limit 3;
-bash: SELECT: command not found
joash_muganda@cloudshell:~/movie-score (movie-prediction-406713)$ bq query --use_legacy_sql=false \
SELECT name, rating, genre, runtime FROM movies.movies_score limit 3;
/usr/lib/google-cloud-sdk/platform/bq/bq.py:17: DeprecationWarning: 'pipes' is deprecated and slated for removal in Python 3.13
import pipes
+-----------------+-----------+--------+---------+
| name | rating | genre | runtime |
+-----------------+-----------+--------+---------+
Love by Drowning	R	Drama	121
It's Just Us	Not Rated	Drama	120
The Robinsons	Not Rated	Action	90
+-----------------+-----------+--------+---------+
joash_muganda@cloudshell:~/movie-score (movie-prediction-406713)$

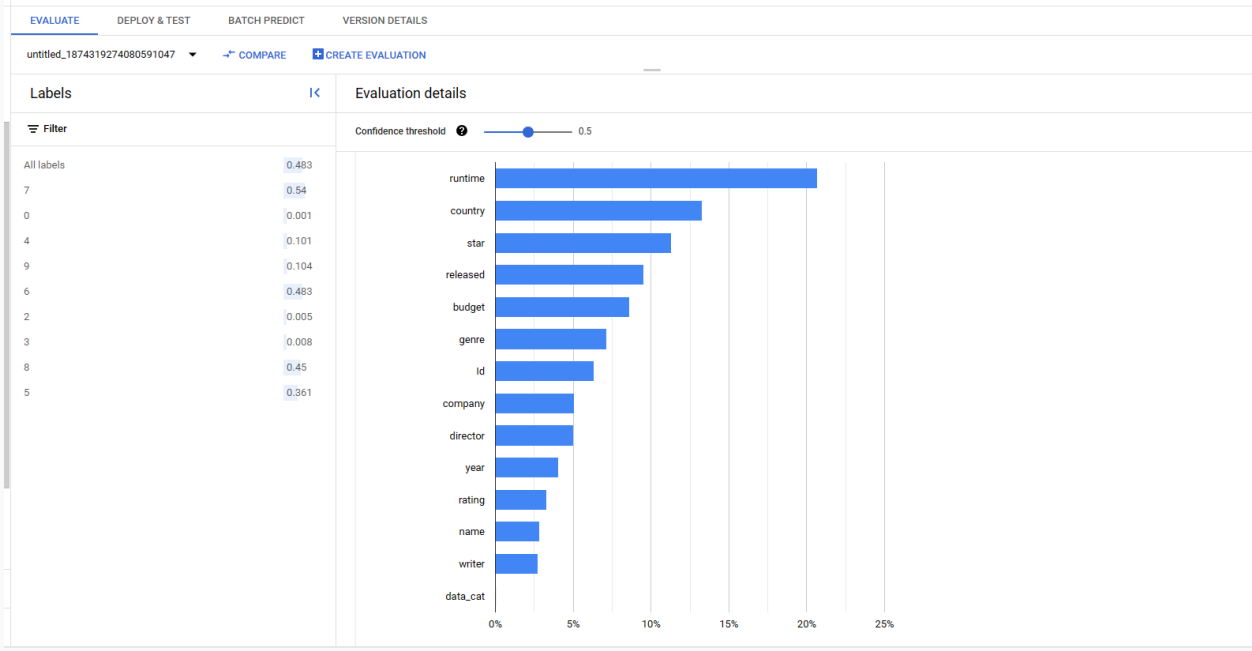
EVALUATE
DEPLOY & TEST
BATCH PREDICT
VERSION DETAILS
untitled_1874319274080591047
COMPARE
CREATE EVALUATION
Labels
Filter
All labels 0.483
7 0.54
0 0.001
4 0.101
9 0.104
6 0.483
2 0.005
3 0.008
8 0.45
5 0.361
Evaluation details
Confidence threshold
0.5
runtime
country
star
released
budget
genre
Id
company
director
year
rating
name
writer
data_cat
0%
5%
10%
15%
20%
25%

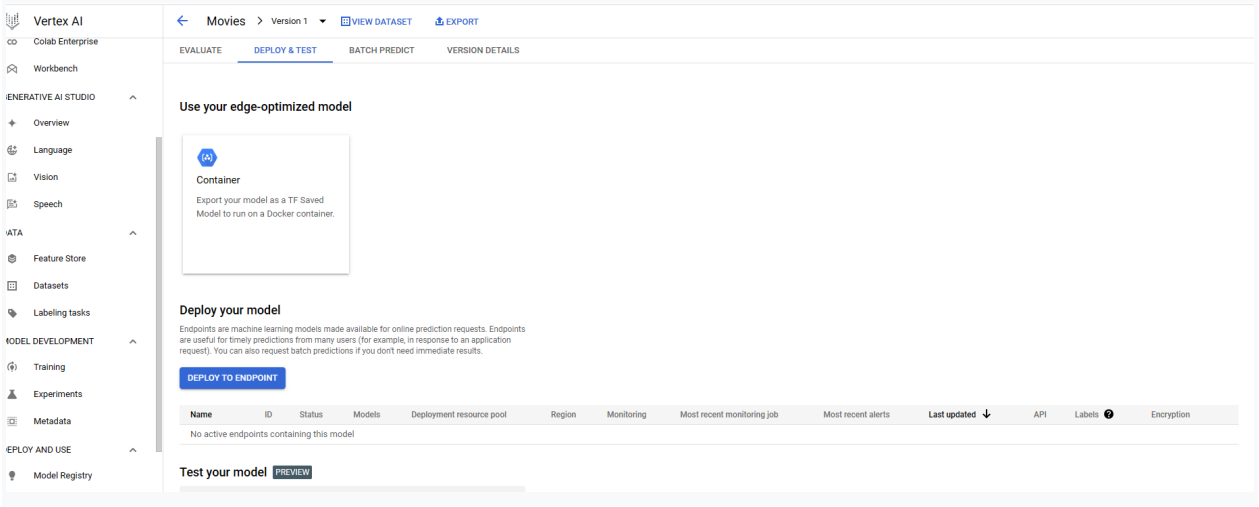
Vertex AI
Colab Enterprise
Workbench
ENERATIVE AI STUDIO
Overview
Language
Vision
Speech
ATA
Feature Store
Datasets
Labeling tasks
ODEL DEVELOPMENT
Training
Experiments
Metadata
EPLOY AND USE
Model Registry
Movies
Version 1
VIEW DATASET
EXPORT
EVALUATE
DEPLOY & TEST
BATCH PREDICT
VERSION DETAILS
Use your edge-optimized model
Container
Export your model as a TF Saved Model to run on a Docker container.
Deploy your model
Endpoints are machine learning models made available for online prediction requests. Endpoints are useful for timely predictions from many users (for example, in response to an application request). You can also request batch predictions if you don't need immediate results.
DEPLOY TO ENDPOINT
Name
ID
Status
Models
Deployment resource pool
Region
Monitoring
Most recent monitoring job
Most recent alerts
Last updated
API
Labels
Encryption
No active endpoints containing this model
Test your model
PREVIEW

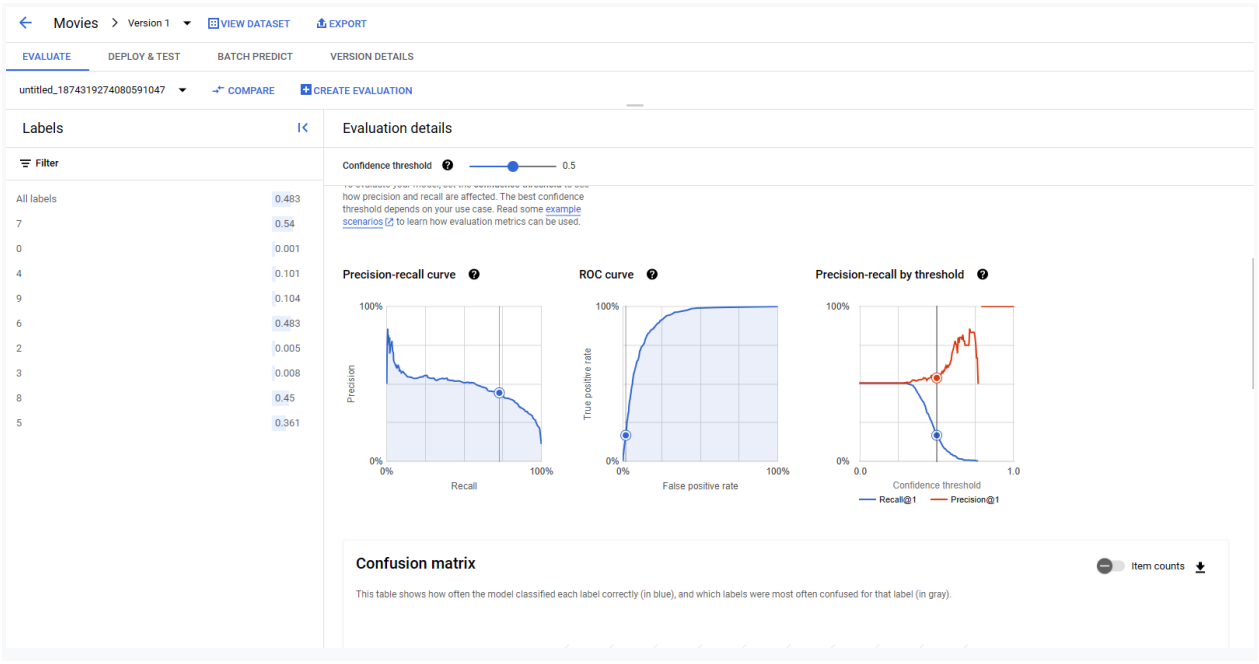
Movies
Version 1
VIEW DATASET
EXPORT
EVALUATE
DEPLOY & TEST
BATCH PREDICT
VERSION DETAILS
untitled_1874319274080591047
COMPARE
CREATE EVALUATION
Labels
Filter
All labels 0.483
7 0.54
0 0.001
4 0.101
9 0.104
6 0.483
2 0.005
3 0.008
8 0.45
5 0.361
Evaluation details
Confidence threshold 0.5
how precision and recall are affected. The best confidence threshold depends on your use case. Read some example scenarios to learn how evaluation metrics can be used.
Precision-recall curve
100%
0%
Precision
Recall
ROC curve
True positive rate
False positive rate
Precision-recall by threshold
0.0
1.0
Confidence threshold
Recall@1
Precision@1
Confusion matrix
Item counts
This table shows how often the model classified each label correctly (in blue), and which labels were most often confused for that label (in gray).

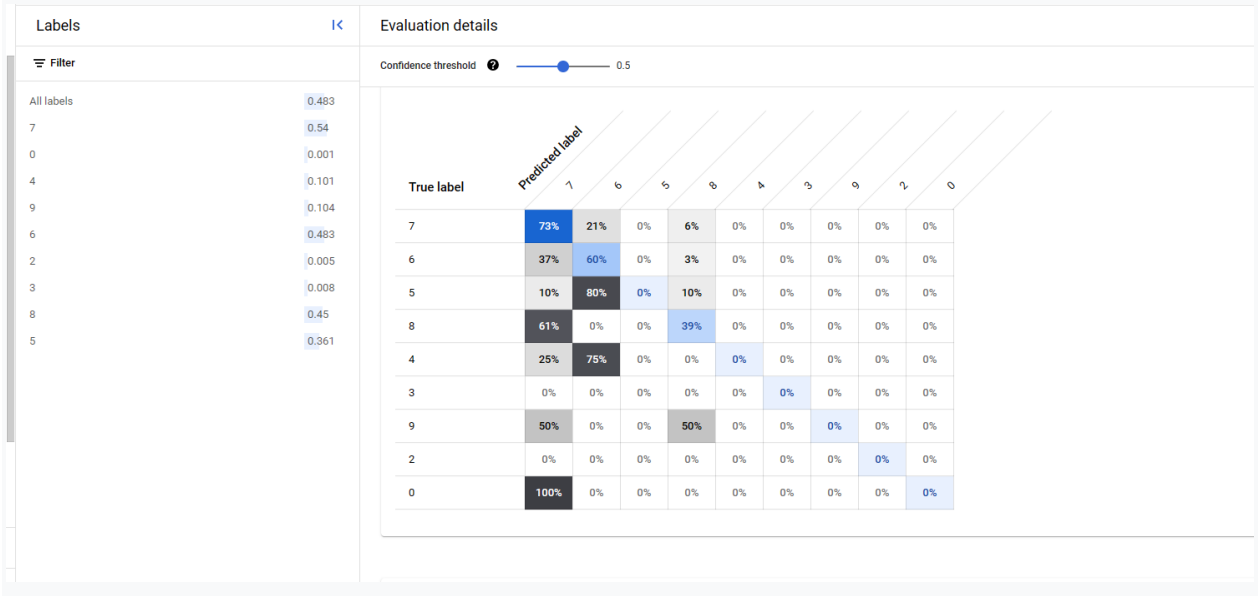
Labels
Filter
All labels 0.483
7 0.54
0 0.001
4 0.101
9 0.104
6 0.483
2 0.005
3 0.008
8 0.45
5 0.361
Evaluation details
Confidence threshold 0.5
Predicted label
True label
7 6 5 8 4 3 9 2 0
7 73% 21% 0% 6% 0% 0% 0% 0% 0%
6 37% 60% 0% 3% 0% 0% 0% 0% 0%
5 10% 80% 0% 10% 0% 0% 0% 0% 0%
8 61% 0% 0% 39% 0% 0% 0% 0% 0%
4 25% 75% 0% 0% 0% 0% 0% 0% 0%
3 0% 0% 0% 0% 0% 0% 0% 0% 0%
9 50% 0% 0% 50% 0% 0% 0% 0% 0%
2 0% 0% 0% 0% 0% 0% 0% 0% 0%
0 100% 0% 0% 0% 0% 0% 0% 0% 0%

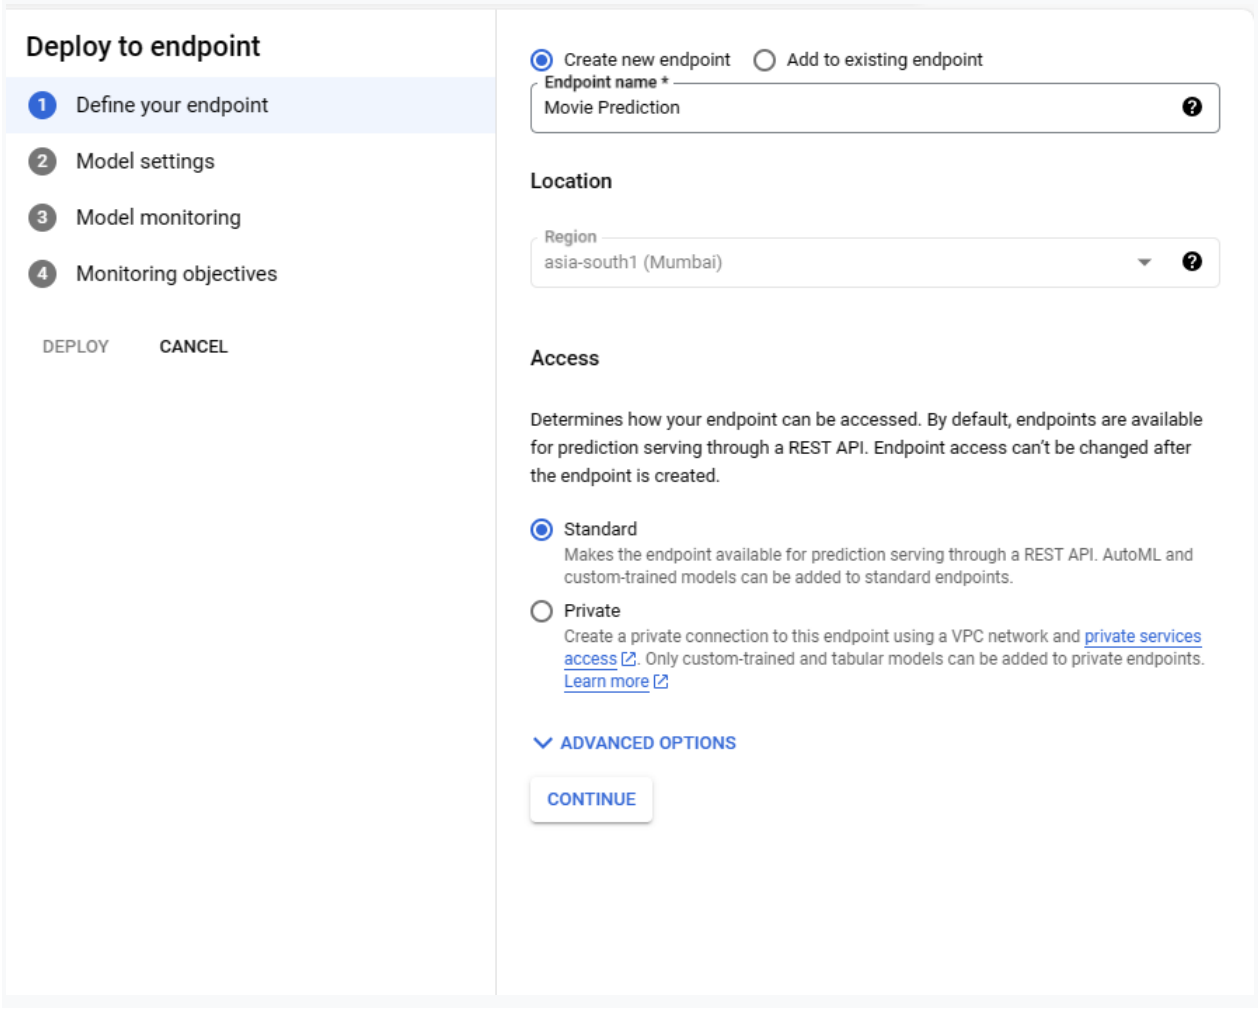
Deploy to endpoint
1 Define your endpoint
2 Model settings
3 Model monitoring
4 Monitoring objectives
DEPLOY
CANCEL
Create new endpoint
Add to existing endpoint
Endpoint name *
Movie Prediction
Location
Region
asia-south1 (Mumbai)
Access
Determines how your endpoint can be accessed. By default, endpoints are available for prediction serving through a REST API. Endpoint access can't be changed after the endpoint is created.
Standard
Makes the endpoint available for prediction serving through a REST API. AutoML and custom-trained models can be added to standard endpoints.
Private
Create a private connection to this endpoint using a VPC network and private services access. Only custom-trained and tabular models can be added to private endpoints. Learn more
ADVANCED OPTIONS
CONTINUE

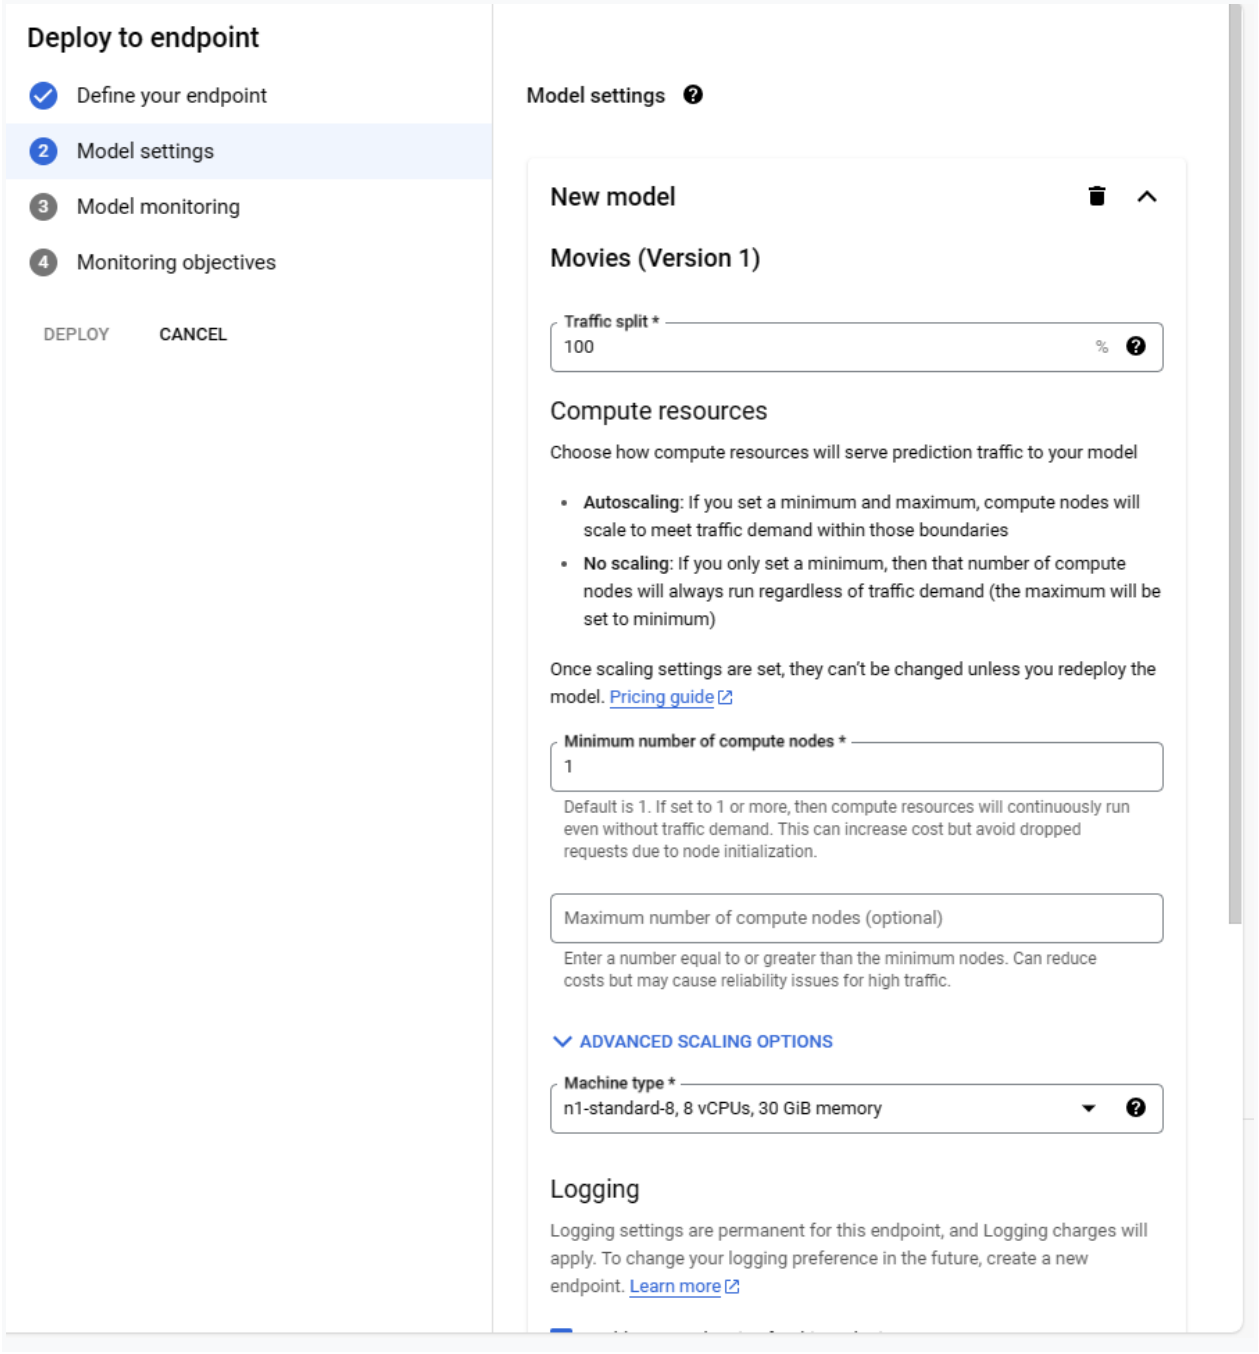
Deploy to endpoint
Define your endpoint
2 Model settings
3 Model monitoring
4 Monitoring objectives
DEPLOY
CANCEL
Model settings
New model
Movies (Version 1)
Traffic split *
100
%
Compute resources
Choose how compute resources will serve prediction traffic to your model
Autoscaling: If you set a minimum and maximum, compute nodes will scale to meet traffic demand within those boundaries
No scaling: If you only set a minimum, then that number of compute nodes will always run regardless of traffic demand (the maximum will be set to minimum)
Once scaling settings are set, they can't be changed unless you redeploy the model. Pricing guide
Minimum number of compute nodes *
1
Default is 1. If set to 1 or more, then compute resources will continuously run even without traffic demand. This can increase cost but avoid dropped requests due to node initialization.
Maximum number of compute nodes (optional)
Enter a number equal to or greater than the minimum nodes. Can reduce costs but may cause reliability issues for high traffic.
ADVANCED SCALING OPTIONS
Machine type *
n1-standard-8, 8 vCPUs, 30 GiB memory
Logging
Logging settings are permanent for this endpoint, and Logging charges will apply. To change your logging preference in the future, create a new endpoint. Learn more

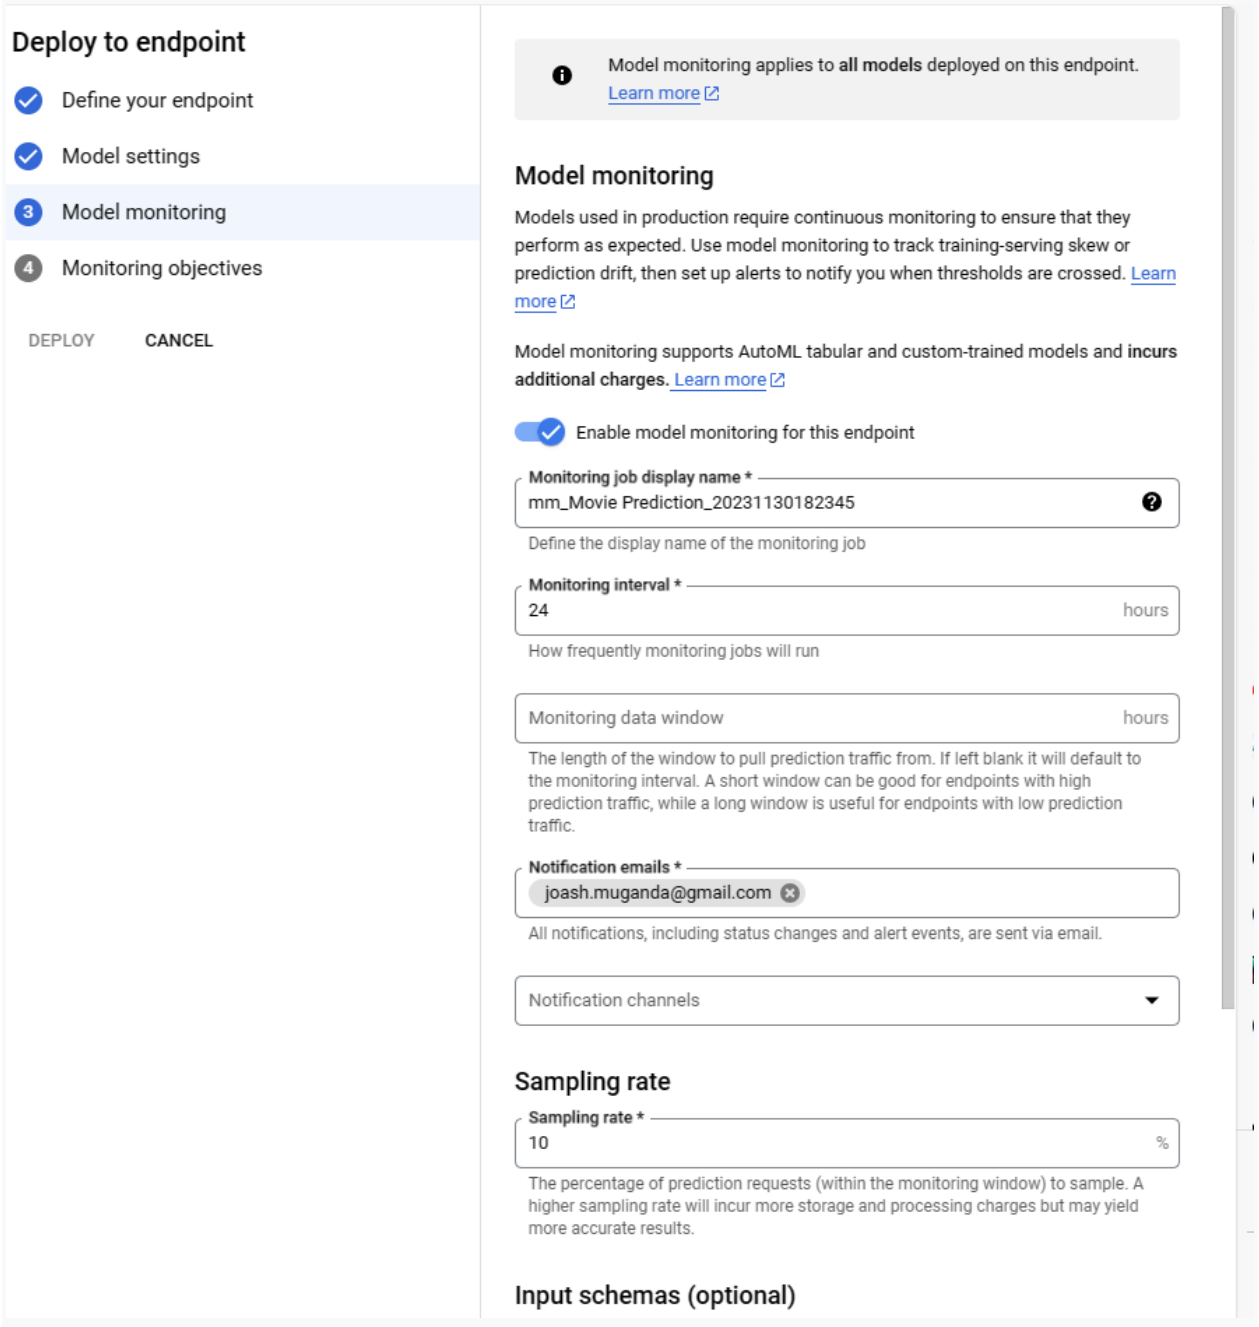
Deploy to endpoint
Define your endpoint
Model settings
3
Model monitoring
4
Monitoring objectives
DEPLOY
CANCEL
Model monitoring applies to all models deployed on this endpoint.
Learn more
Model monitoring
Models used in production require continuous monitoring to ensure that they perform as expected. Use model monitoring to track training-serving skew or prediction drift, then set up alerts to notify you when thresholds are crossed. Learn more
Model monitoring supports AutoML tabular and custom-trained models and incurs additional charges. Learn more
Enable model monitoring for this endpoint
Monitoring job display name *
mm_Movie Prediction_20231130182345
Define the display name of the monitoring job
Monitoring interval *
24
hours
How frequently monitoring jobs will run
Monitoring data window
hours
The length of the window to pull prediction traffic from. If left blank it will default to the monitoring interval. A short window can be good for endpoints with high prediction traffic, while a long window is useful for endpoints with low prediction traffic.
Notification emails *
joash.muganda@gmail.com
All notifications, including status changes and alert events, are sent via email.
Notification channels
Sampling rate
Sampling rate *
10
%
The percentage of prediction requests (within the monitoring window) to sample. A higher sampling rate will incur more storage and processing charges but may yield more accurate results.
Input schemas (optional)

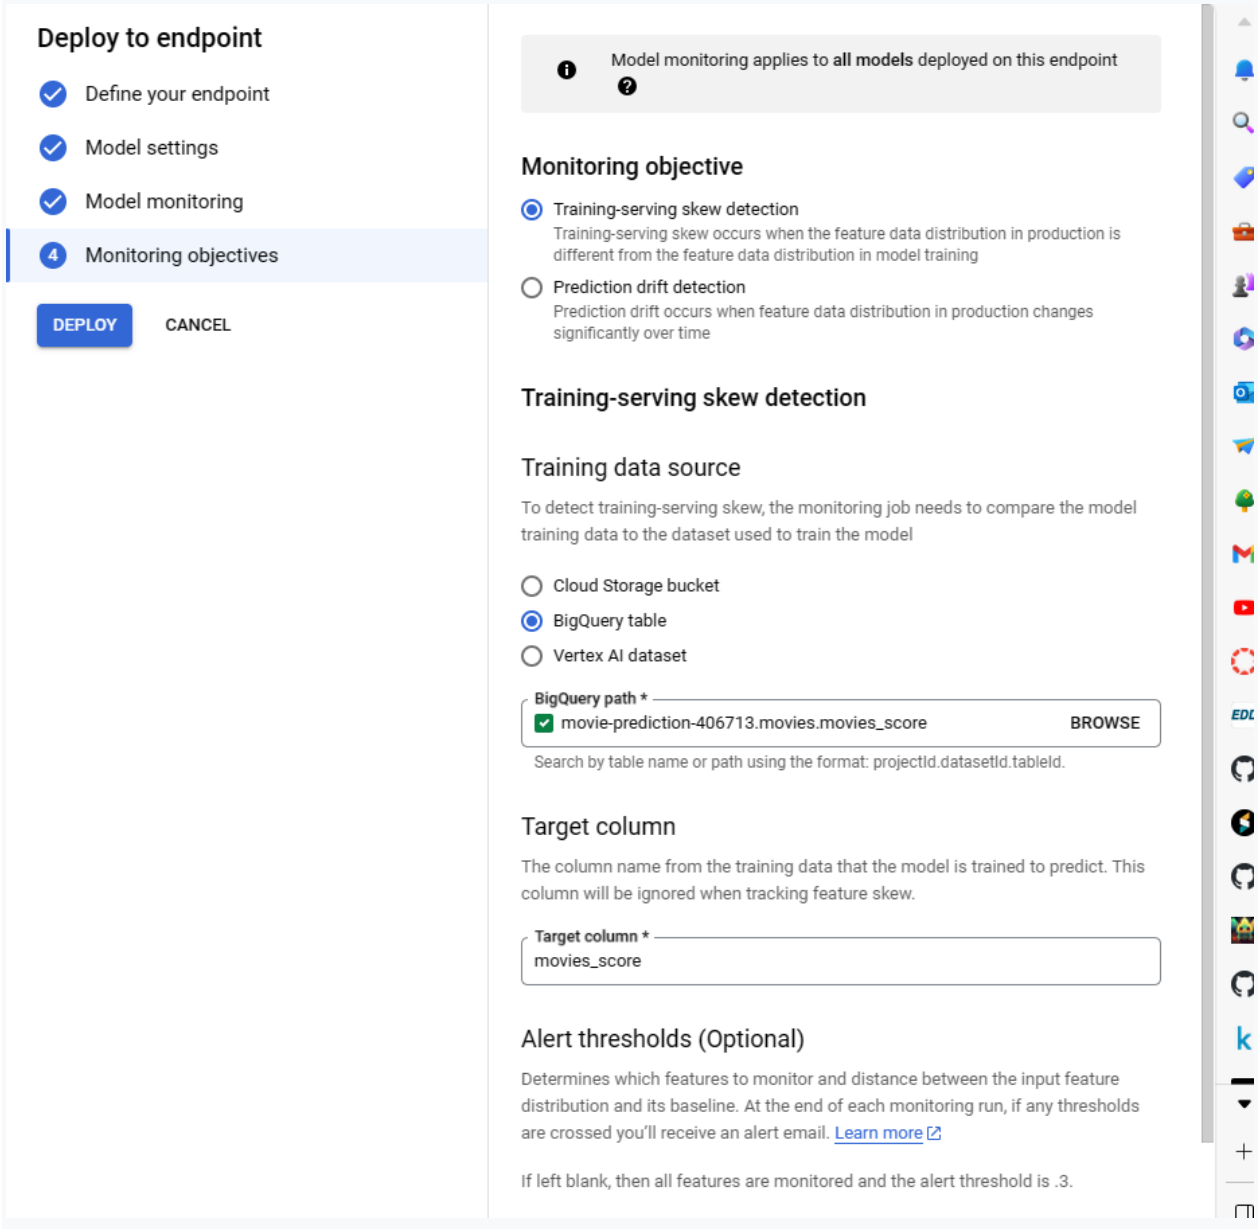
Deploy to endpoint
Define your endpoint
Model settings
Model monitoring
4
Monitoring objectives
DEPLOY
CANCEL
Model monitoring applies to all models deployed on this endpoint
Monitoring objective
Training-serving skew detection
Training-serving skew occurs when the feature data distribution in production is different from the feature data distribution in model training
Prediction drift detection
Prediction drift occurs when feature data distribution in production changes significantly over time
Training-serving skew detection
Training data source
To detect training-serving skew, the monitoring job needs to compare the model training data to the dataset used to train the model
Cloud Storage bucket
BigQuery table
Vertex AI dataset
BigQuery path *
movie-prediction-406713.movies.movies_score
BROWSE
Search by table name or path using the format: projectId.datasetId.tableId.
Target column
The column name from the training data that the model is trained to predict. This column will be ignored when tracking feature skew.
Target column *
movies_score
Alert thresholds (Optional)
Determines which features to monitor and distance between the input feature distribution and its baseline. At the end of each monitoring run, if any thresholds are crossed you'll receive an alert email. Learn more
If left blank, then all features are monitored and the alert threshold is .3.

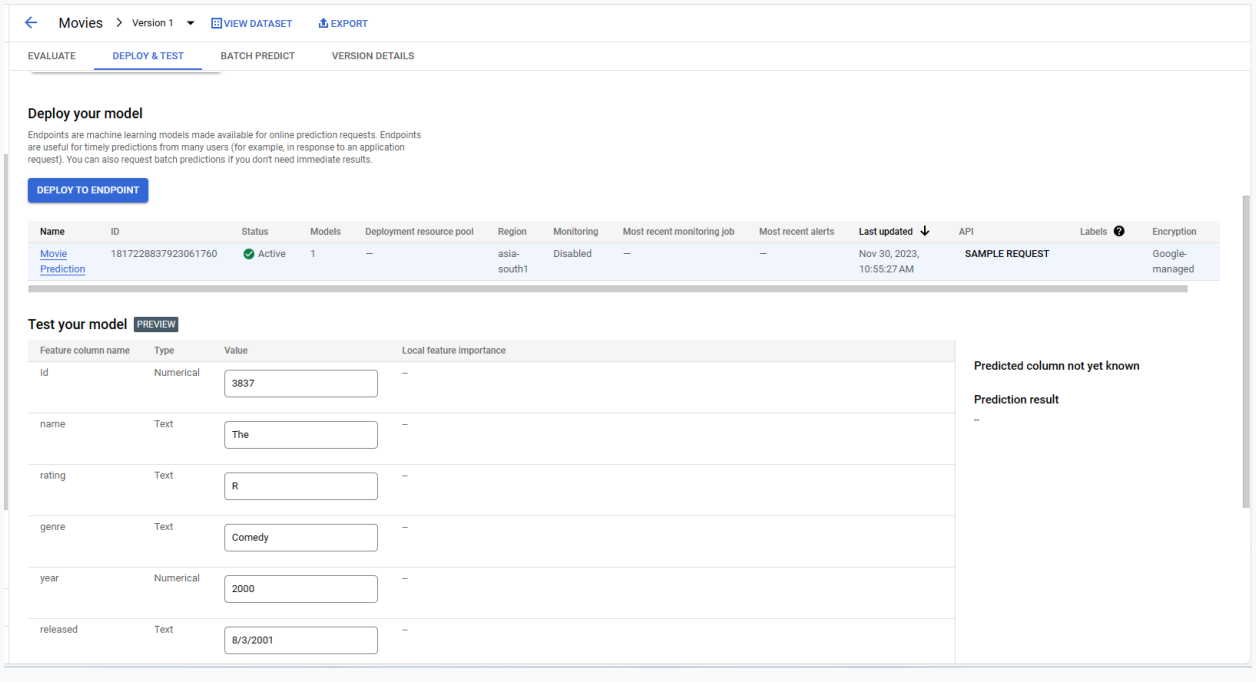
Movies
Version 1
VIEW DATASET
EXPORT
EVALUATE
DEPLOY & TEST
BATCH PREDICT
VERSION DETAILS
Deploy your model
Endpoints are machine learning models made available for online prediction requests. Endpoints are useful for timely predictions from many users (for example, in response to an application request). You can also request batch predictions if you don't need immediate results.
DEPLOY TO ENDPOINT
Name
ID
Status
Models
Deployment resource pool
Region
Monitoring
Most recent monitoring job
Most recent alerts
Last updated
API
Labels
Encryption
Movie Prediction
1817228837923061760
Active
1
–
asia-south1
Disabled
–
–
Nov 30, 2023, 10:55:27 AM
SAMPLE REQUEST
Google-managed
Test your model
PREVIEW
Feature column name
Type
Value
Local feature importance
Id
Numerical
3837
name
Text
The
rating
R
genre
Comedy
year
2000
released
8/3/2001
Predicted column not yet known
Prediction result

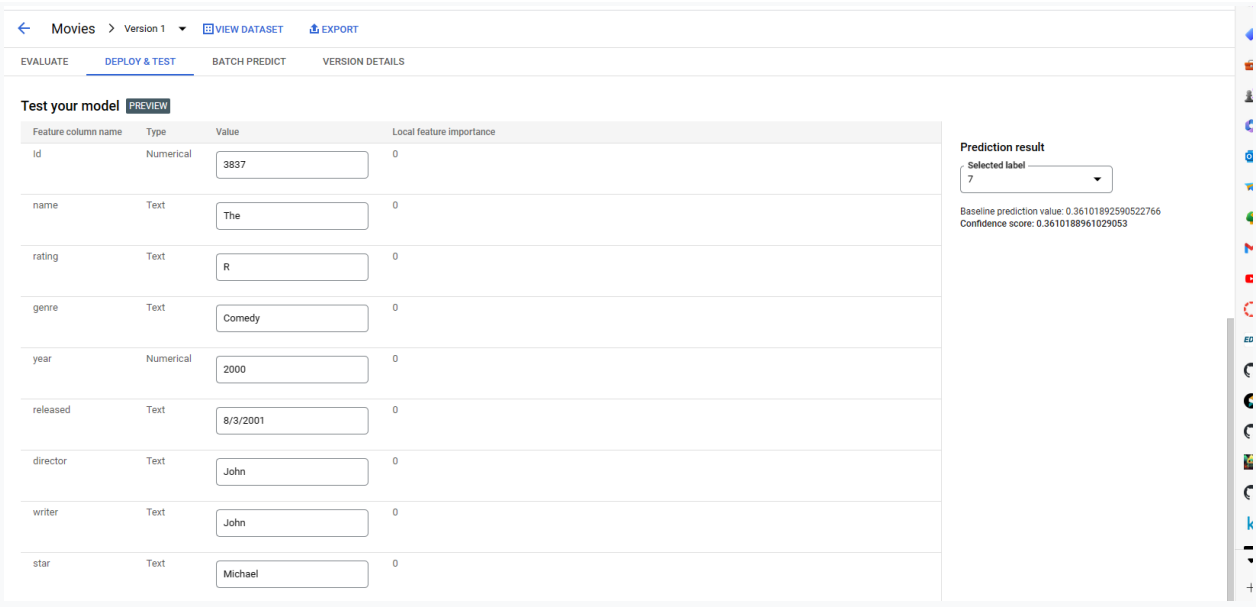
Movies
Version 1
VIEW DATASET
EXPORT
EVALUATE
DEPLOY & TEST
BATCH PREDICT
VERSION DETAILS
Test your model
PREVIEW
Feature column name
Type
Value
Local feature importance
Id
Numerical
3837
name
Text
The
rating
R
genre
Comedy
year
2000
released
8/3/2001
director
John
writer
John
star
Michael
Prediction result
Selected label
7
Baseline prediction value: 0.36101892590522766
Confidence score: 0.3610188961029053

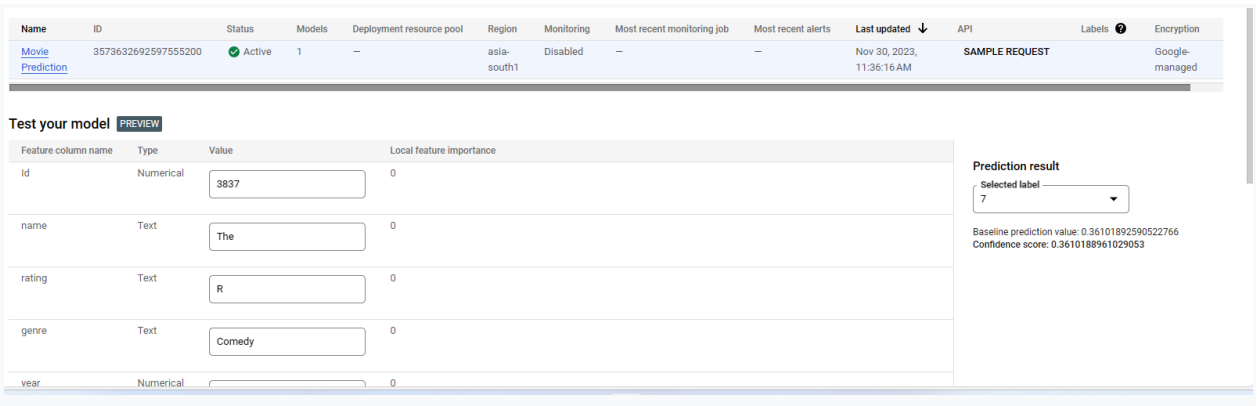
Name
ID
Status
Models
Deployment resource pool
Region
Monitoring
Most recent monitoring job
Most recent alerts
Last updated
API
Labels
Encryption
Movie Prediction
3573632692597555200
Active
1
–
asia-south1
Disabled
–
–
Nov 30, 2023, 11:36:16 AM
SAMPLE REQUEST
Google-managed
Test your model
PREVIEW
Feature column name
Type
Value
Local feature importance
Id
Numerical
3837
name
Text
The
rating
R
genre
Comedy
year
Numerical
Prediction result
Selected label
7
Baseline prediction value: 0.36101892590522766
Confidence score: 0.3610188961029053

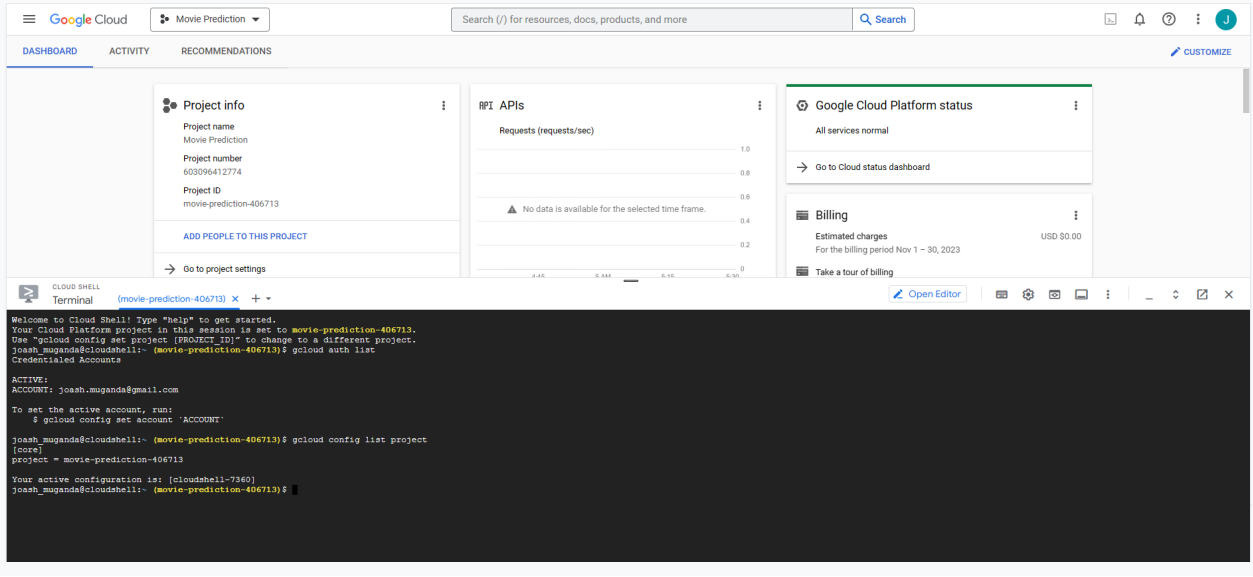
Google Cloud
Movie Prediction
Search (/) for resources, docs, products, and more
Search
DASHBOARD
ACTIVITY
RECOMMENDATIONS
CUSTOMIZE
Project info
Project name
Movie Prediction
Project number
603096412774
Project ID
movie-prediction-406713
ADD PEOPLE TO THIS PROJECT
Go to project settings
APIs
Requests (requests/sec)
No data is available for the selected time frame.
Google Cloud Platform status
All services normal
Go to Cloud status dashboard
Billing
Estimated charges
USD $0.00
For the billing period Nov 1 – 30, 2023
Take a tour of billing
CLOUD SHELL
Terminal
(movie-prediction-406713)
Open Editor
Welcome to Cloud Shell! Type "help" to get started.
Your Cloud Platform project in this session is set to movie-prediction-406713.
Use "gcloud config set project [PROJECT_ID]" to change to a different project.
joash_muganda@cloudshell:~ (movie-prediction-406713)$ gcloud auth list
Credentialed Accounts
ACTIVE:
ACCOUNT: joash.muganda@gmail.com
To set the active account, run:
$ gcloud config set account 'ACCOUNT'
joash_muganda@cloudshell:~ (movie-prediction-406713)$ gcloud config list project
[core]
project = movie-prediction-406713
Your active configuration is: [cloudshell-7360]
joash_muganda@cloudshell:~ (movie-prediction-406713)$

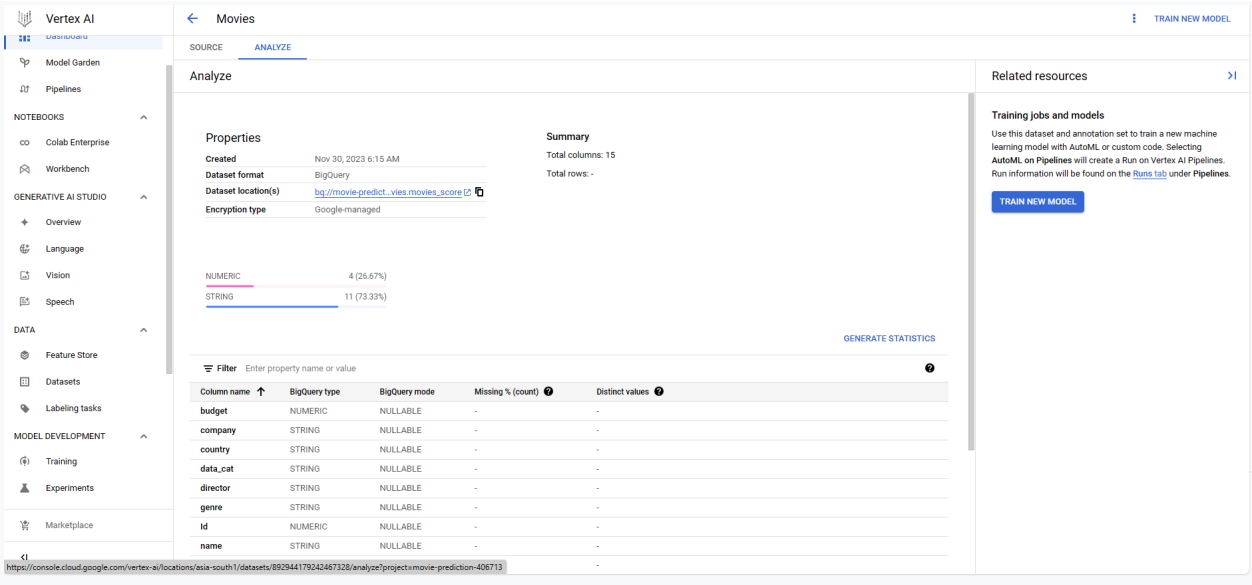
Vertex AI
Model Garden
Pipelines
NOTEBOOKS
Colab Enterprise
Workbench
GENERATIVE AI STUDIO
Overview
Language
Vision
Speech
DATA
Feature Store
Datasets
Labeling tasks
MODEL DEVELOPMENT
Training
Experiments
Marketplace
Movies
TRAIN NEW MODEL
SOURCE
ANALYZE
Analyze
Properties
Created
Nov 30, 2023 6:15 AM
Dataset format
BigQuery
Dataset location(s)
bq://movie-predict...vies.movies_score
Encryption type
Google-managed
Summary
Total columns: 15
Total rows: -
NUMERIC
4 (26.67%)
STRING
11 (73.33%)
GENERATE STATISTICS
Filter
Enter property name or value
Column name
BigQuery type
BigQuery mode
Missing % (count)
Distinct values
budget NUMERIC NULLABLE
company STRING NULLABLE
country STRING NULLABLE
data_cat STRING NULLABLE
director STRING NULLABLE
genre STRING NULLABLE
id NUMERIC NULLABLE
name STRING NULLABLE
Related resources
Training jobs and models
Use this dataset and annotation set to train a new machine learning model with AutoML or custom code. Selecting AutoML on Pipelines will create a Run on Vertex AI Pipelines. Run information will be found on the Runs tab under Pipelines.
TRAIN NEW MODEL
https://console.cloud.google.com/vertex-ai/locations/asia-south1/datasets/8929441792424673 28/analyze?project=movie-prediction-406713

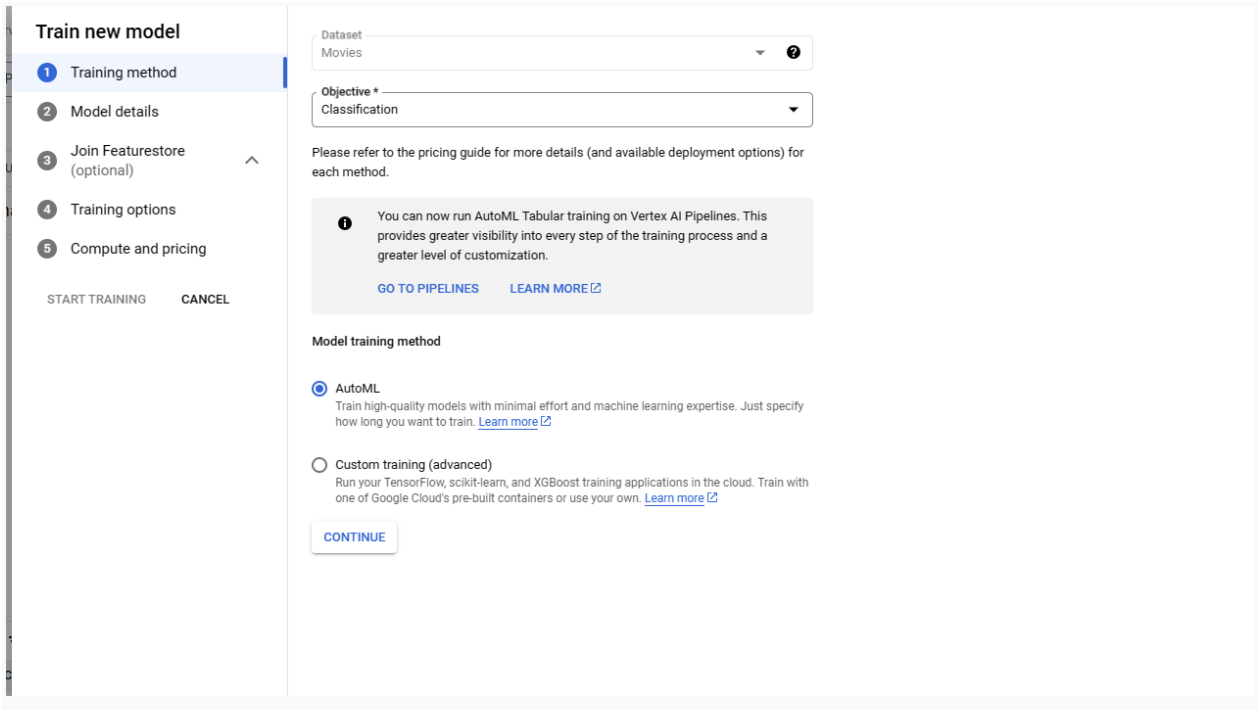
Train new model
1 Training method
2 Model details
3 Join Featurestore (optional)
4 Training options
5 Compute and pricing
START TRAINING
CANCEL
Dataset
Movies
Objective *
Classification
Please refer to the pricing guide for more details (and available deployment options) for each method.
You can now run AutoML Tabular training on Vertex AI Pipelines. This provides greater visibility into every step of the training process and a greater level of customization.
GO TO PIPELINES
LEARN MORE
Model training method
AutoML
Train high-quality models with minimal effort and machine learning expertise. Just specify how long you want to train. Learn more
Custom training (advanced)
Run your TensorFlow, scikit-learn, and XGBoost training applications in the cloud. Train with one of Google Cloud's pre-built containers or use your own. Learn more
CONTINUE

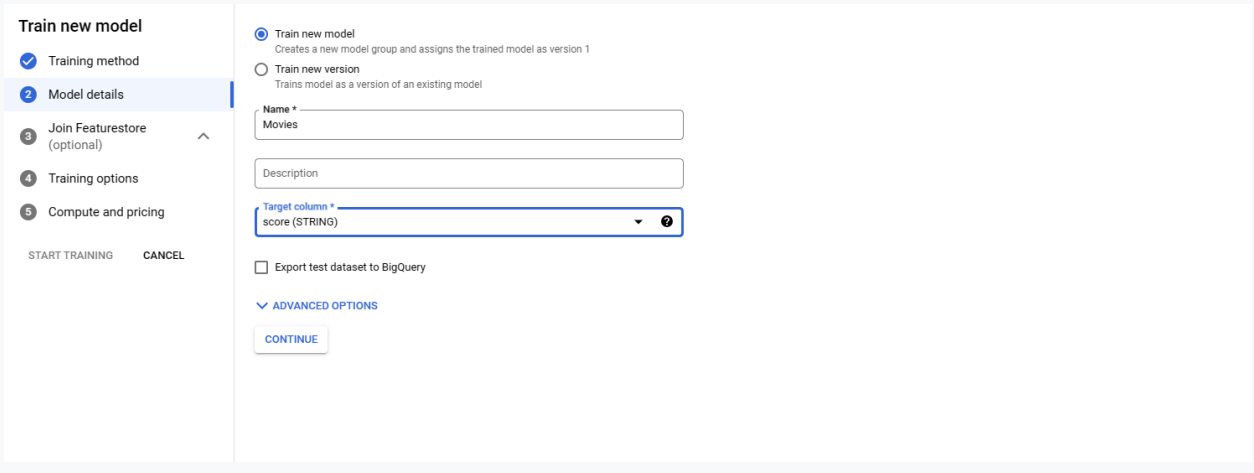

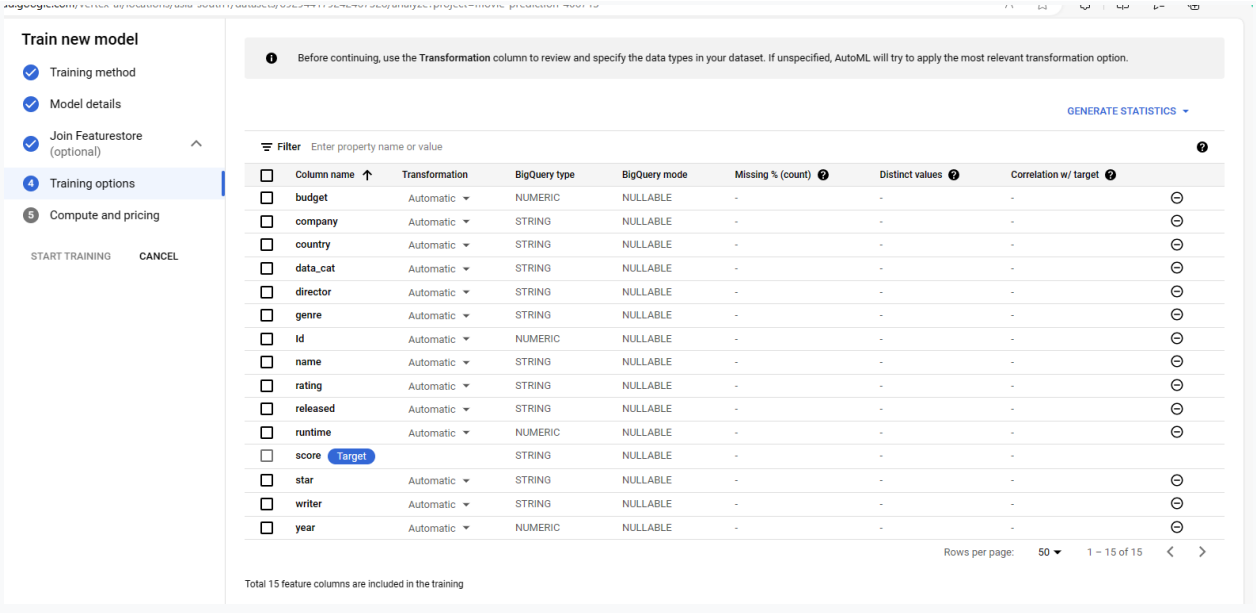

2. Predicting Daily MelbourneTemperatures

Note: I wasn't permitted to use the Iowa Liquor sales dataset. Used the Melbourne temperatures dataset instead.

Training Model:

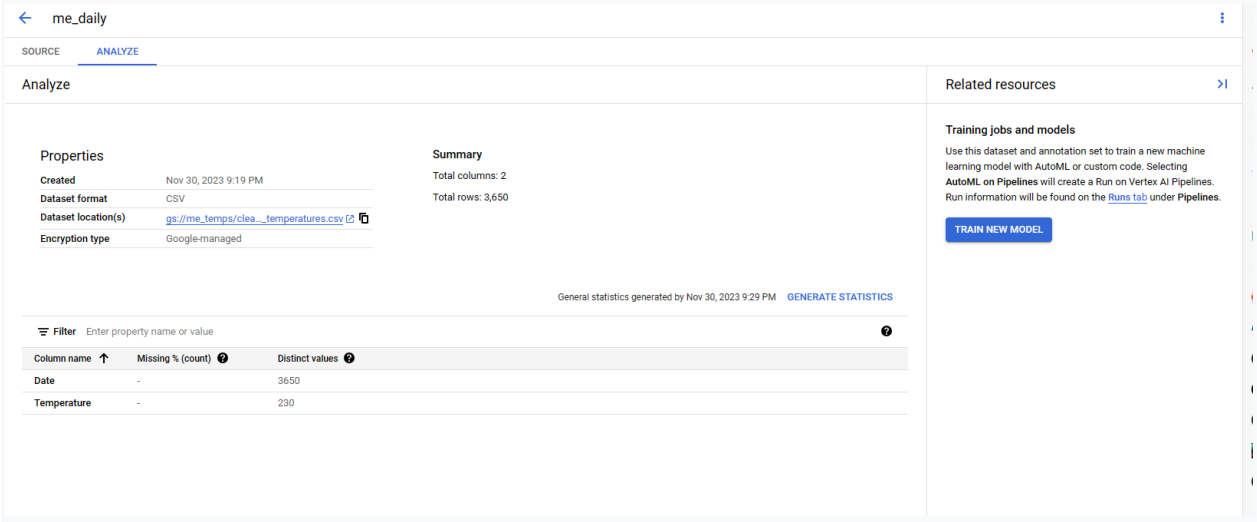
me_daily
SOURCE
ANALYZE
Analyze
Related resources
Properties
Created
Nov 30, 2023 9:19 PM
Dataset format
CSV
Dataset location(s)
gs://me_temps/clea..._temperatures.csv
Encryption type
Google-managed
Summary
Total columns: 2
Total rows: 3,650
Training jobs and models
Use this dataset and annotation set to train a new machine learning model with AutoML or custom code. Selecting AutoML on Pipelines will create a Run on Vertex AI Pipelines. Run information will be found on the Runs tab under Pipelines.
TRAIN NEW MODEL
General statistics generated by Nov 30, 2023 9:29 PM
GENERATE STATISTICS
Filter
Enter property name or value
Column name
Missing % (count)
Distinct values
Date
-
3650
Temperature
-
230

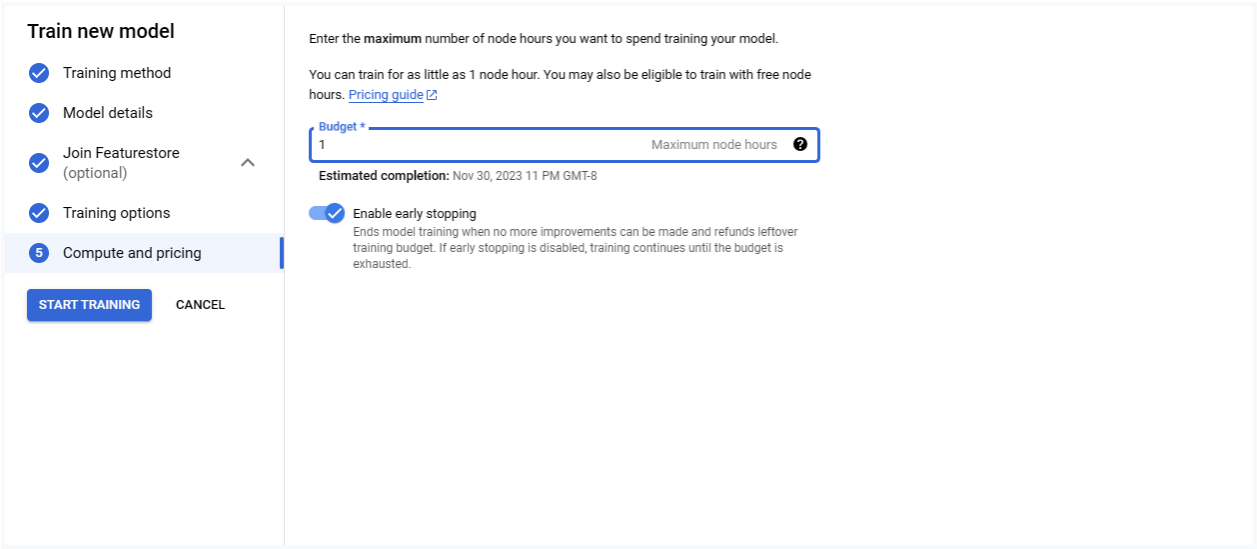
Train new model
Training method
Model details
Join Featurestore (optional)
Training options
5
Compute and pricing
START TRAINING
CANCEL
Enter the maximum number of node hours you want to spend training your model.
You can train for as little as 1 node hour. You may also be eligible to train with free node hours.
Pricing guide
Budget *
1
Maximum node hours
Estimated completion: Nov 30, 2023 11 PM GMT-8
Enable early stopping
Ends model training when no more improvements can be made and refunds leftover training budget. If early stopping is disabled, training continues until the budget is exhausted.

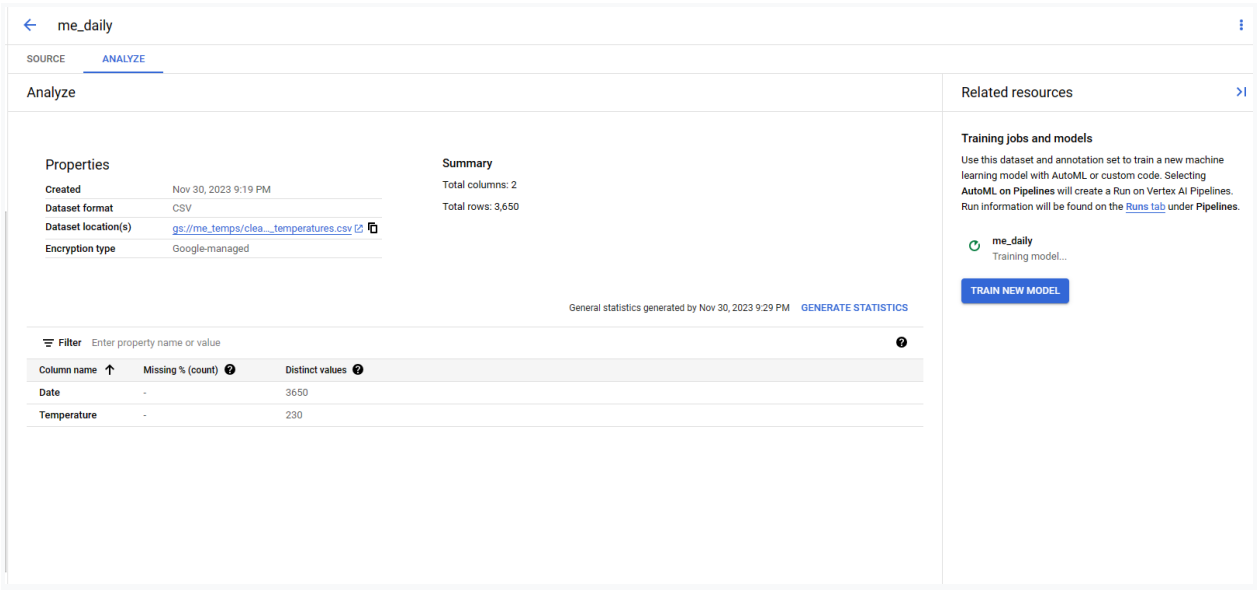
me_daily
SOURCE
ANALYZE
Analyze
Related resources
Properties
Created
Nov 30, 2023 9:19 PM
Dataset format
CSV
Dataset location(s)
gs://me_temps/clea..._temperatures.csv
Encryption type
Google-managed
Summary
Total columns: 2
Total rows: 3,650
Training jobs and models
Use this dataset and annotation set to train a new machine learning model with AutoML or custom code. Selecting AutoML on Pipelines will create a Run on Vertex AI Pipelines. Run information will be found on the Runs tab under Pipelines.
me_daily
Training model...
TRAIN NEW MODEL
General statistics generated by Nov 30, 2023 9:29 PM
GENERATE STATISTICS
Filter
Enter property name or value
Column name
Missing % (count)
Distinct values
Date
-
3650
Temperature
-
230

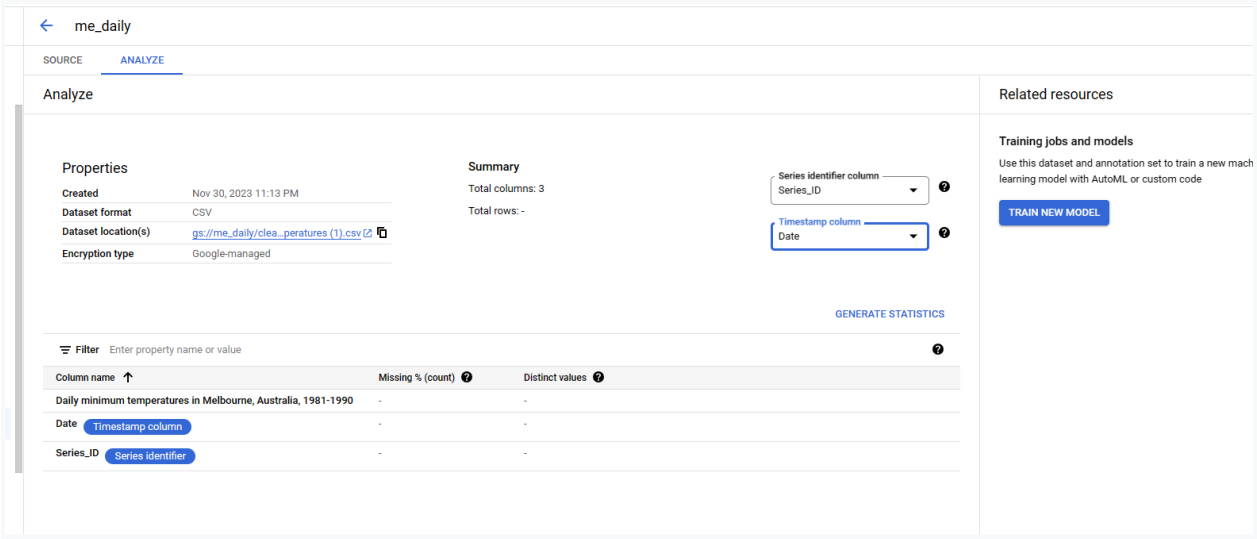
me_daily
SOURCE
ANALYZE
Analyze
Properties
Created
Nov 30, 2023 11:13 PM
Dataset format
CSV
Dataset location(s)
gs://me_daily/clea...peratures (1).csv
Encryption type
Google-managed
Summary
Total columns: 3
Total rows: -
Series identifier column
Series_ID
Timestamp column
Date
GENERATE STATISTICS
Filter
Enter property name or value
Column name
Missing % (count)
Distinct values
Daily minimum temperatures in Melbourne, Australia, 1981-1990
Date
Timestamp column
Series_ID
Series identifier
Related resources
Training jobs and models
Use this dataset and annotation set to train a new mach
learning model with AutoML or custom code
TRAIN NEW MODEL

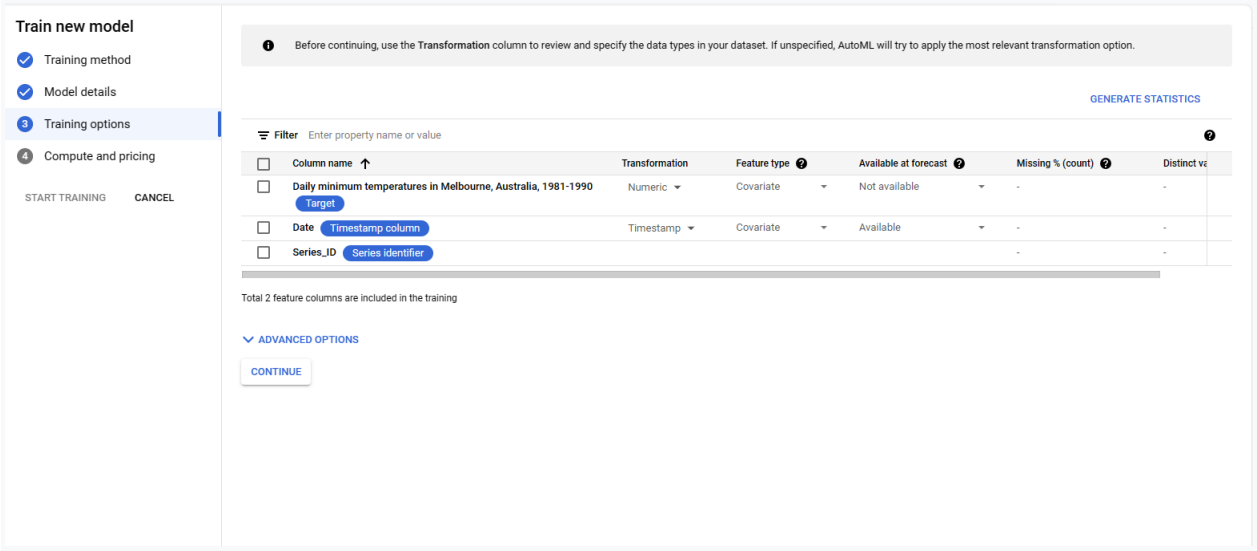
Train new model
Training method
Model details
Training options
Compute and pricing
START TRAINING
CANCEL
Before continuing, use the Transformation column to review and specify the data types in your dataset. If unspecified, AutoML will try to apply the most relevant transformation option.
GENERATE STATISTICS
Filter
Enter property name or value
Column name
Transformation
Feature type
Available at forecast
Missing % (count)
Distinct va
Daily minimum temperatures in Melbourne, Australia, 1981-1990
Target
Numeric
Covariate
Not available
Date
Timestamp column
Timestamp
Covariate
Available
Series_ID
Series identifier
Total 2 feature columns are included in the training
ADVANCED OPTIONS
CONTINUE

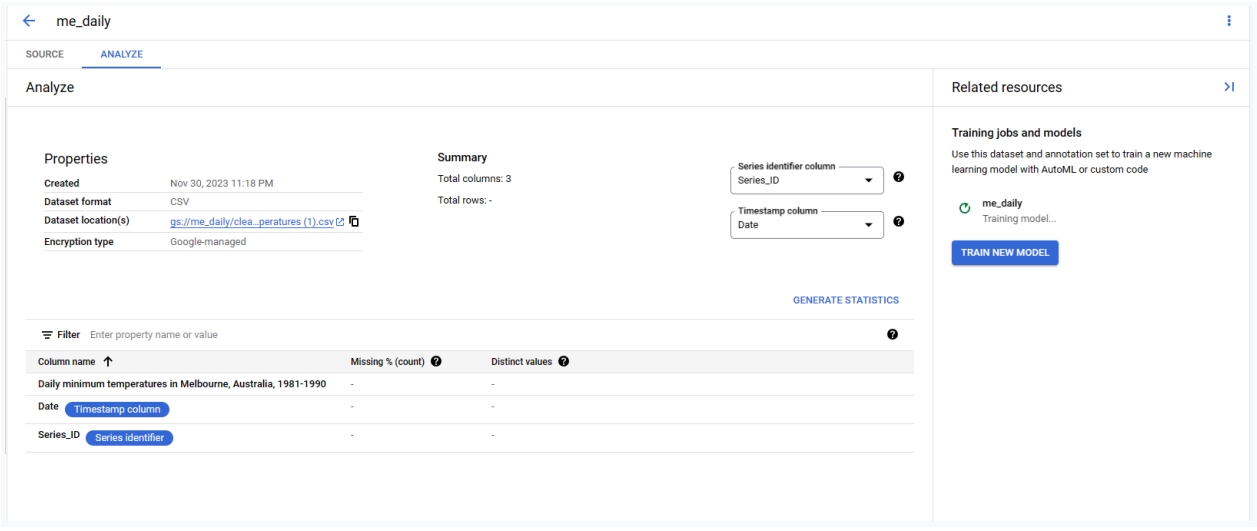
me_daily
SOURCE
ANALYZE
Analyze
Properties
Created
Nov 30, 2023 11:18 PM
Dataset format
CSV
Dataset location(s)
gs://me_daily/clea...peratures (1).csv
Encryption type
Google-managed
Summary
Total columns: 3
Total rows: -
Series identifier column
Series_ID
Timestamp column
Date
GENERATE STATISTICS
Filter
Enter property name or value
Column name
Missing % (count)
Distinct values
Daily minimum temperatures in Melbourne, Australia, 1981-1990
Date
Timestamp column
Series_ID
Series identifier
Related resources
Training jobs and models
Use this dataset and annotation set to train a new machine
learning model with AutoML or custom code
me_daily
Training model...
TRAIN NEW MODEL

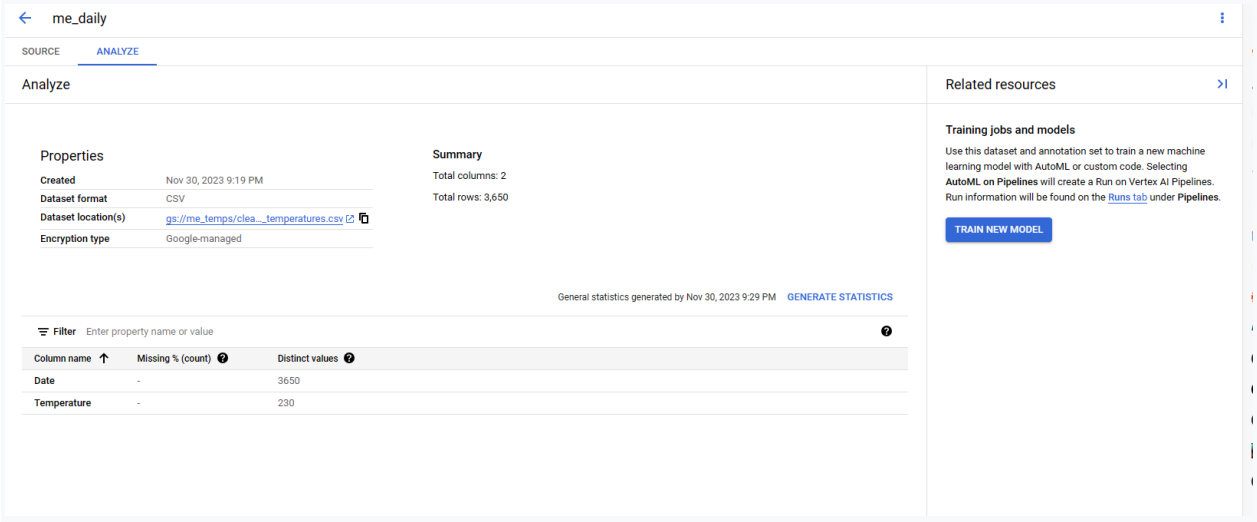
me_daily
SOURCE
ANALYZE
Analyze
Related resources
Properties
Created
Nov 30, 2023 9:19 PM
Dataset format
CSV
Dataset location(s)
gs://me_temps/clea..._temperatures.csv
Encryption type
Google-managed
Summary
Total columns: 2
Total rows: 3,650
Training jobs and models
Use this dataset and annotation set to train a new machine learning model with AutoML or custom code. Selecting AutoML on Pipelines will create a Run on Vertex AI Pipelines. Run information will be found on the Runs tab under Pipelines.
TRAIN NEW MODEL
General statistics generated by Nov 30, 2023 9:29 PM
GENERATE STATISTICS
Filter
Enter property name or value
Column name
Missing % (count)
Distinct values
Date
-
3650
Temperature
-
230

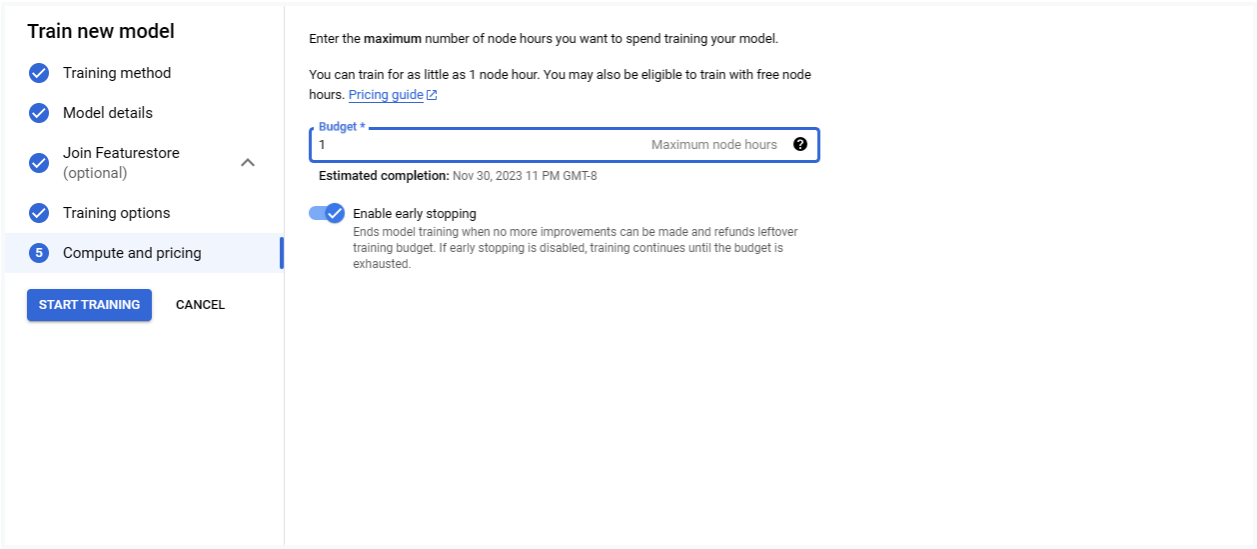
Train new model
Training method
Model details
Join Featurestore (optional)
Training options
5
Compute and pricing
START TRAINING
CANCEL
Enter the maximum number of node hours you want to spend training your model.
You can train for as little as 1 node hour. You may also be eligible to train with free node hours.
Pricing guide
Budget *
1
Maximum node hours
Estimated completion: Nov 30, 2023 11 PM GMT-8
Enable early stopping
Ends model training when no more improvements can be made and refunds leftover training budget. If early stopping is disabled, training continues until the budget is exhausted.

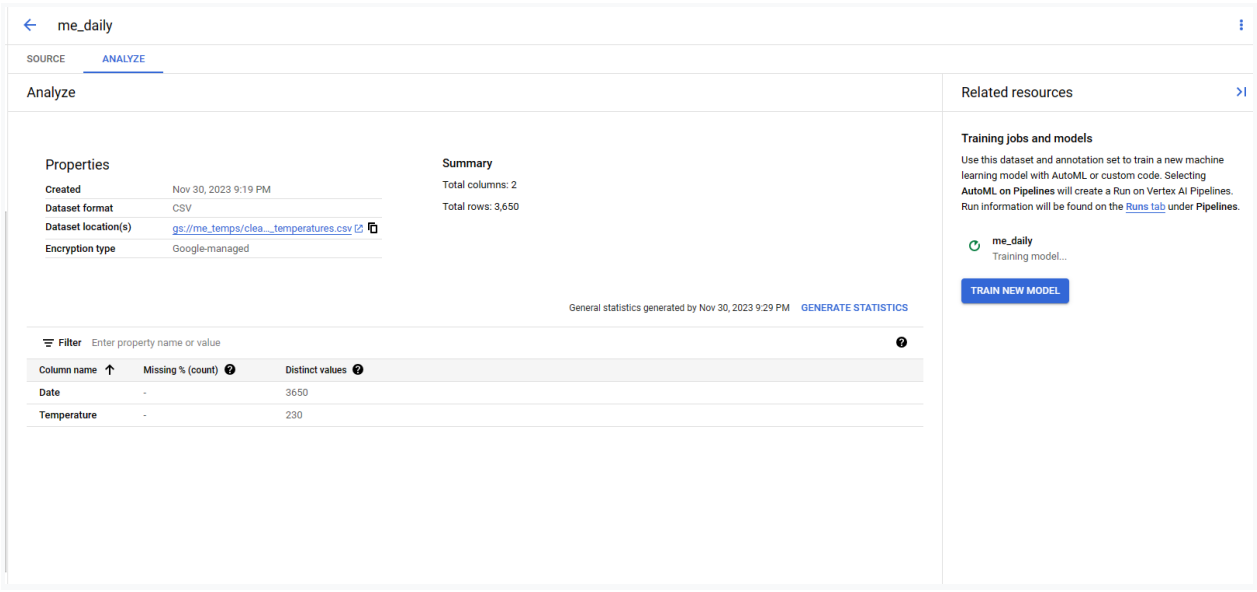
me_daily
SOURCE
ANALYZE
Analyze
Related resources
Properties
Created
Nov 30, 2023 9:19 PM
Dataset format
CSV
Dataset location(s)
gs://me_temps/clea..._temperatures.csv
Encryption type
Google-managed
Summary
Total columns: 2
Total rows: 3,650
Training jobs and models
Use this dataset and annotation set to train a new machine learning model with AutoML or custom code. Selecting AutoML on Pipelines will create a Run on Vertex AI Pipelines. Run information will be found on the Runs tab under Pipelines.
me_daily
Training model...
TRAIN NEW MODEL
General statistics generated by Nov 30, 2023 9:29 PM
GENERATE STATISTICS
Filter
Enter property name or value
Column name
Missing % (count)
Distinct values
Date
-
3650
Temperature
-
230

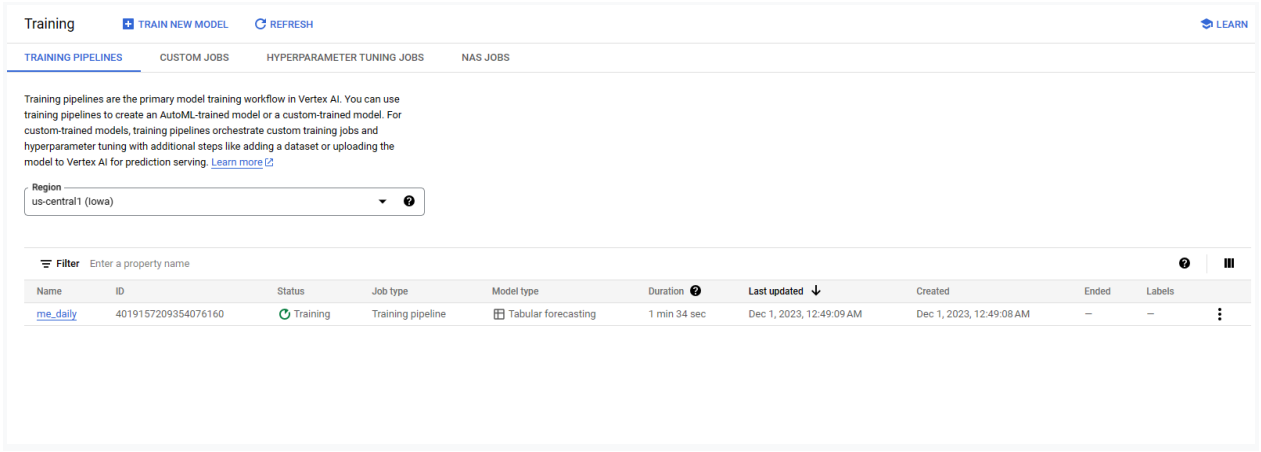

Trained Mode:

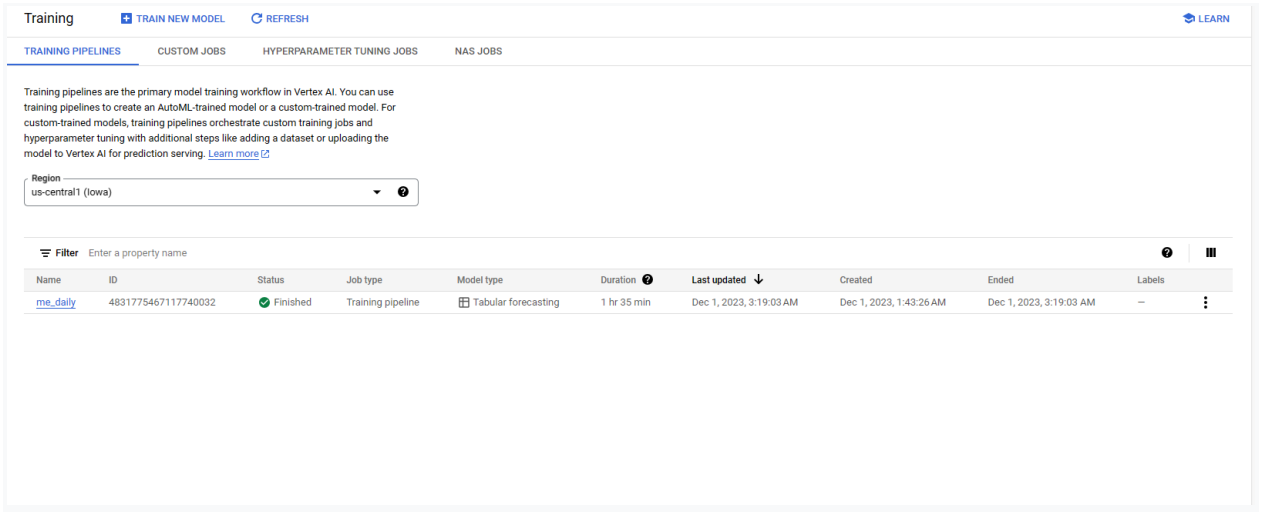

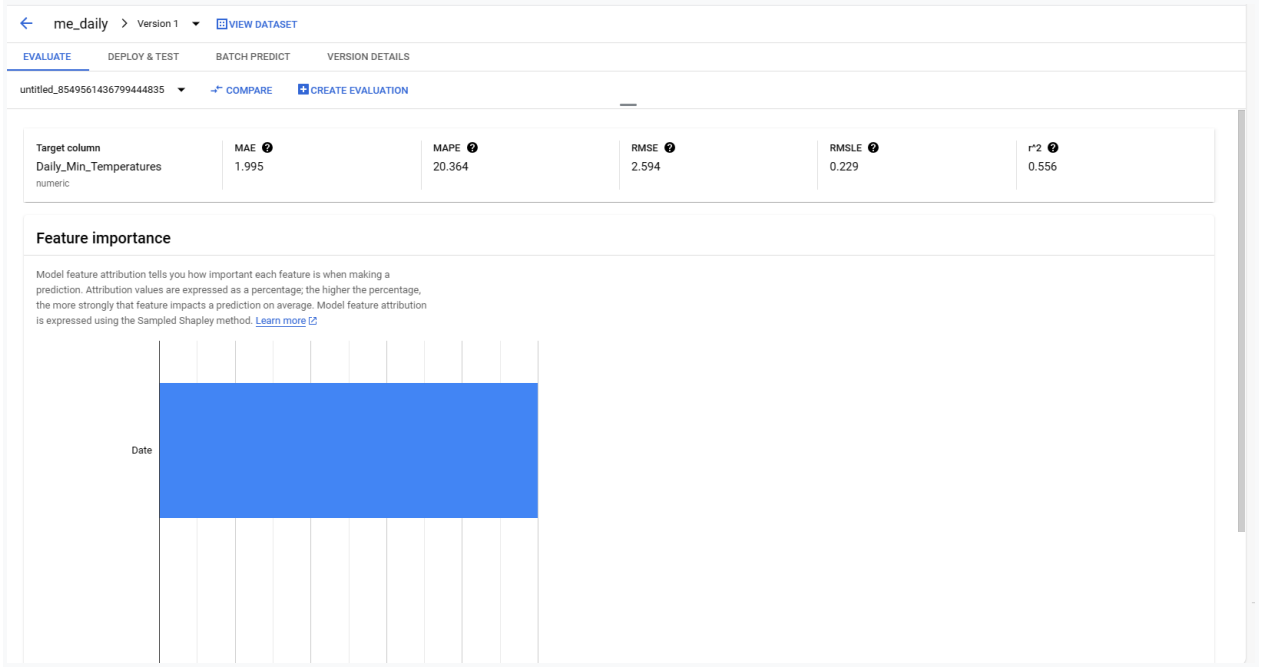

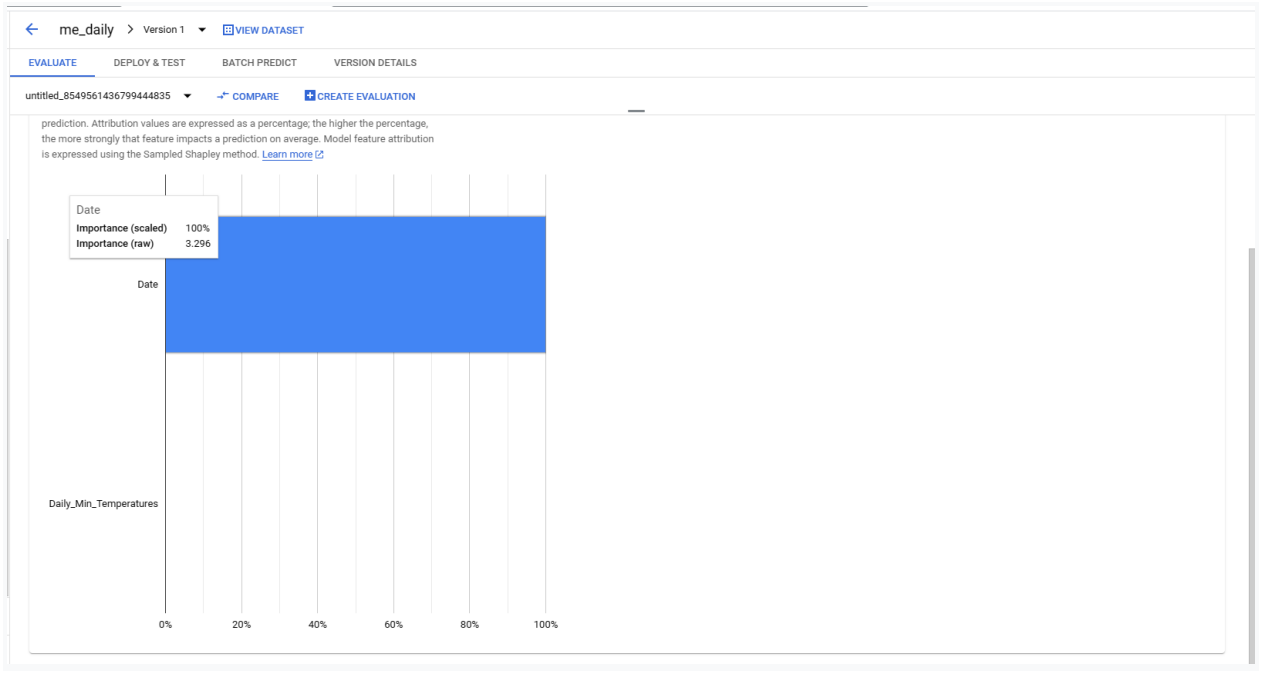
me_daily
Version 1
VIEW DATASET
EVALUATE
DEPLOY & TEST
BATCH PREDICT
VERSION DETAILS
untitled_8549561436799444835
COMPARE
CREATE EVALUATION
prediction. Attribution values are expressed as a percentage; the higher the percentage, the more strongly that feature impacts a prediction on average. Model feature attribution is expressed using the Sampled Shapley method. Learn more
Date
Importance (scaled) 100%
Importance (raw) 3.296
Date
Daily_Min_Temperatures
0%
20%
40%
60%
80%
100%

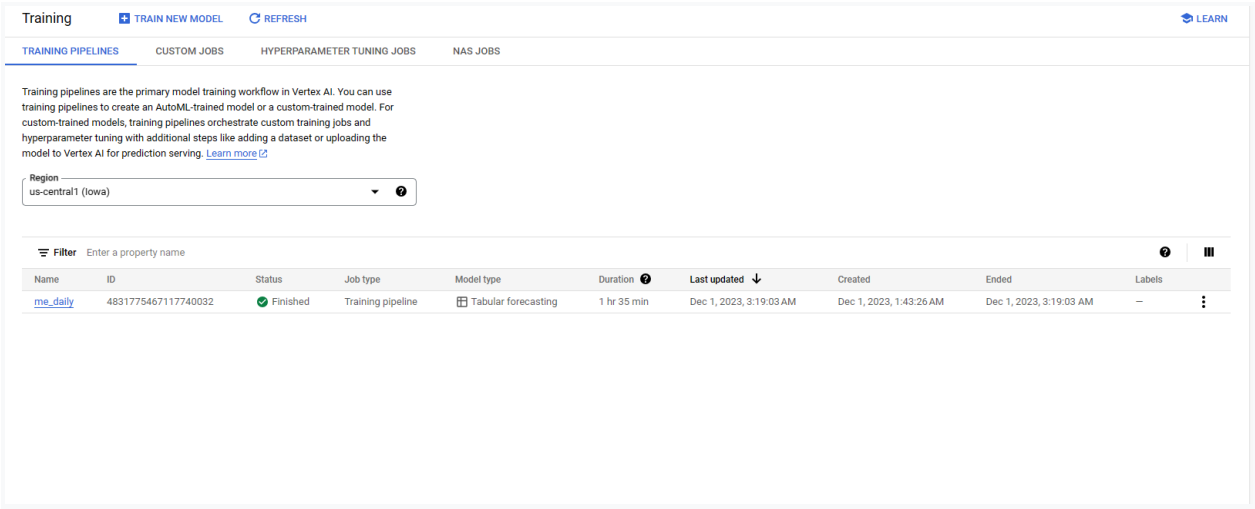

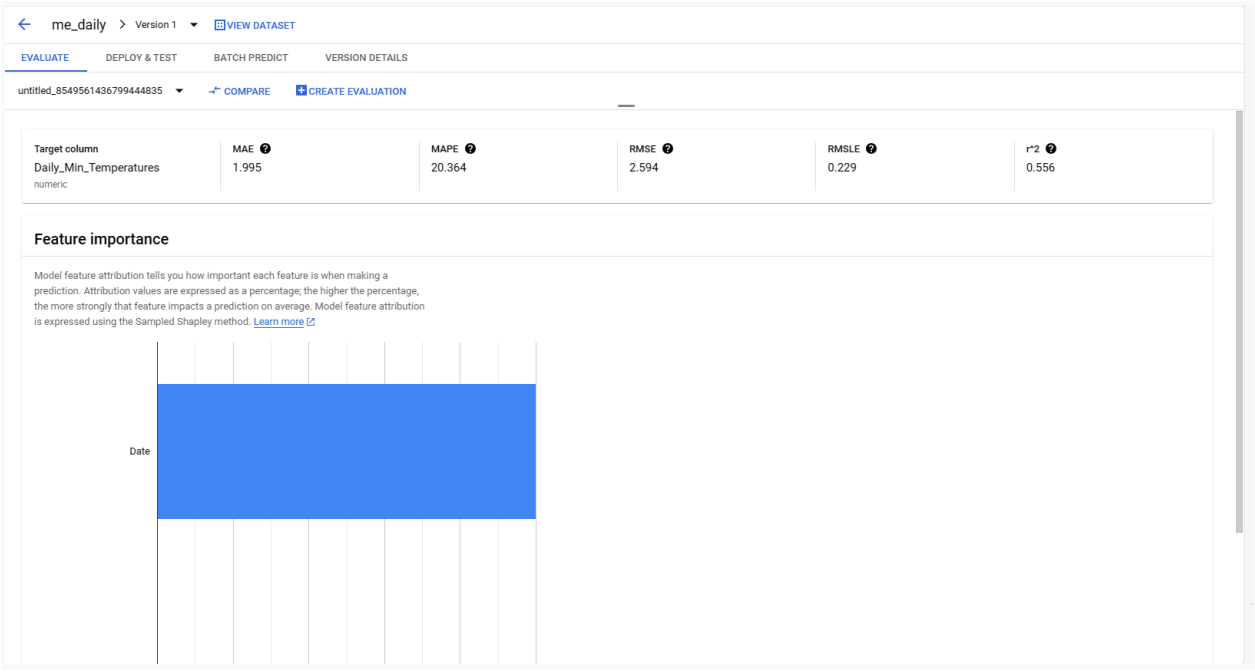

Deployment Failed:

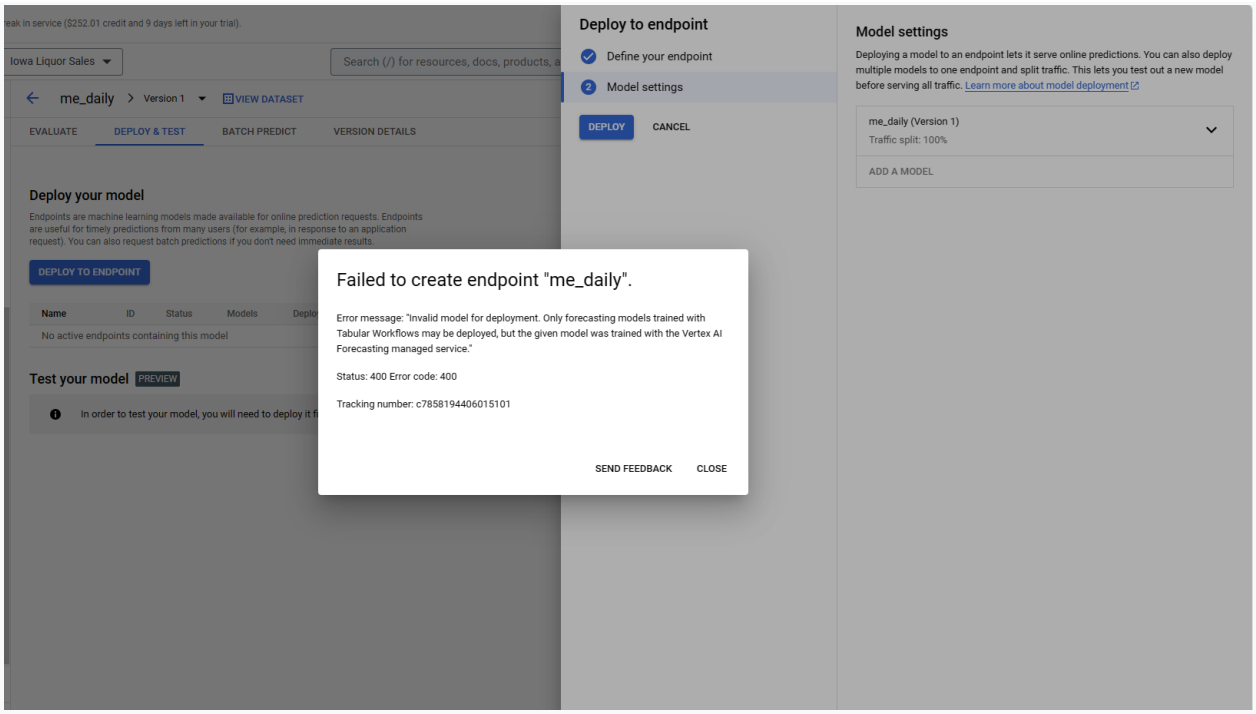
Iowa Liquor Sales
Search (/) for resources, docs, products,
me_daily
Version 1
VIEW DATASET
EVALUATE
DEPLOY & TEST
BATCH PREDICT
VERSION DETAILS
Deploy your model
Endpoints are machine learning models made available for online prediction requests. Endpoints are useful for timely predictions from many users (for example, in response to an application request). You can also request batch predictions if you don't need immediate results.
DEPLOY TO ENDPOINT
Name
ID
Status
Models
No active endpoints containing this model
Test your model
PREVIEW
In order to test your model, you will need to deploy it
Failed to create endpoint "me_daily".
Error message: "Invalid model for deployment. Only forecasting models trained with Tabular Workflows may be deployed, but the given model was trained with the Vertex AI Forecasting managed service."
Status: 400 Error code: 400
Tracking number: c7858194406015101
SEND FEEDBACK
CLOSE
Deploy to endpoint
Define your endpoint
2
Model settings
DEPLOY
CANCEL
Model settings
Deploying a model to an endpoint lets it serve online predictions. You can also deploy multiple models to one endpoint and split traffic. This lets you test out a new model before serving all traffic. Learn more about model deployment
me_daily (Version 1)
Traffic split: 100%
ADD A MODEL